

GhostBed

AI Visibility Audit

An external audit of how AI engines and search see GhostBed today. And what we would do about it.

Prepared for	GhostBed leadership
Date	June 22, 2026
Scope	Public site, search and AI surface, reputation, agent-commerce readiness
Contact	ethanpzhou@berkeley.edu
Prepared by	Armaan Priyadarshan, Samarjit Deshmukh, Ethan Zhou, Shamit D'Souza
Dartmouth Math & CS · Dartmouth CS · UC Berkeley MET (EECS & Business) · Stanford CS & Biology	

AI visibility scorecard

We graded GhostBed on five dimensions of how AI engines and search treat the brand. Each dimension is scored 0 to 100 against observed behavior across AI answer engines, the category listicles those engines cite, third-party reputation sources, the brand's own indexed pages, and the agent-commerce surface. Competitor scores are our estimates from the same public signals. Composite is the simple average.



GhostBed has the best agent-commerce plumbing in its competitive set and one of the cleaner product-schema layers, yet it is the set's least visible brand in the AI answers that sell mattresses. Transactability, at 84, leads Saatva, Purple, and Nectar, none of which expose a live agent endpoint. The two dimensions that decide whether an engine ever names GhostBed to a new shopper, recommendability (38) and reputation (48), are the lowest. The rails are ready. The engines never route a buyer onto them.

DIMENSION	GHOSTBED	SAATVA	PURPLE	NECTAR	WHAT IT MEASURES
Discoverability	64	74	80	64	Entity presence, agent discovery files, crawler access, third-party listings
Quotability	70	66	74	62	Schema depth, structured product and policy data, accuracy in machine-readable form
Recommendability	38	90	76	82	Inclusion in AI category recommendations and the listicles engines cite
Transactability	84	46	46	46	UCP and MCP endpoints, payment handlers, agent-readable pricing
Reputation	48	84	70	68	Review aggregator presence and rating, sentiment, links from the entity to that reputation
Composite	61 (C)	74 (C)	70 (C)	64 (C)	GhostBed trails the set average by about 9 points, all of it on recommendability and reputation, while leading every competitor on the agent-commerce layer.

Grading scale: A 90+, B 75 to 89, C 60 to 74, D 50 to 59, F below 50. Competitor scores are estimates from observed public signals, not a full audit of those brands. Method and queries in the methodology section.

The short version

GhostBed is an established cooling-mattress brand: a decade in market, a Venus Williams signature line, Costco and Amazon distribution, and a Shopify store that already runs the modern agent-commerce rails most brands have not built. The product itself reviews well. The problem is upstream of all of it. When we asked AI answer engines for the best cooling mattress, GhostBed's own category, it was not named once. Saatva, Nectar, Purple, and Helix were. The checkout an AI agent would use is live and correct. The engine never sends a buyer to it.

Six findings, ordered worst first. The sixth is a strength.

1. **AI engines do not recommend GhostBed, not even for cooling.** Asked for the best cooling mattress, the best online mattress, and the best mattress overall, AI answer engines named Saatva, Nectar, Purple, Helix, DreamCloud, and WinkBed. GhostBed appeared in none of the three, despite "Luxury Cooling Mattresses" being its own homepage title.
2. **The reputation a machine can read is both weak and unlinked.** GhostBed's product pages advertise a 4.8 from 10,223 reviews, its own widget. The open web shows 3.5 on Trustpilot (82% one-star all time), 1.61 at the BBB, and 1.5 on ConsumerAffairs. The Organization in GhostBed's schema carries sameAs: none, so nothing connects the brand to any of it.
3. **The homepage is strong for humans and blank for machines.** Good title, meta, and social cards, but zero JSON-LD on the homepage: no Organization, WebSite, or Brand entity, and no server-rendered h1. Purple ships all of it.
4. **GhostBed wrote the content AI engines quote, then hid it from them.** 36 head-to-head comparison pages (GhostBed vs Saatva, vs Casper, vs Helix) and roughly 86 buying guides, with no Article or FAQPage schema, and an llms.txt that by its own text only "mirrors" the agent file rather than mapping any of it.
5. **Product schema is strong; the entity and policy schema around it is missing.** PDPs render full Product, Offer, and AggregateRating. There is no WebSite, BreadcrumbList, or FAQPage, so an engine can read a price but not the trial or warranty terms.
6. **The agent-commerce layer is live, and ahead of the entire category.** A working UCP profile, an MCP endpoint, agents.md, and an agentic-discovery sitemap. Saatva, Purple, and Nectar expose none of it. This is the asset to build on.

ChatGPT, Perplexity, and Google AI Mode all default to live web retrieval in 2026. Every mattress query that hits live retrieval reads the same category listicles and the same reputation corpus, and reaches Saatva, Nectar, Purple, and Helix. GhostBed built the checkout rails for that exact moment. It is just never the brand the engine routes the buyer to.

AI engines do not recommend GhostBed, not even for cooling

AI answer engines are now where many mattress shoppers start. They do not browse; they ask "what is the best cooling mattress" and get back a short, synthesized list. We ran three framings of that question through web-search-backed retrieval, the same retrieval ChatGPT, Perplexity, and Google AI Mode use by default. GhostBed was absent from all three.

Query, web-search-backed AI answer: "best cooling mattress for hot sleepers 2026"

"*Helix Midnight Luxe (best overall), Brooklyn Bedding Aurora Luxe, WinkBed, DreamCloud Premier Hybrid, Bear Elite...*"

GhostBed is not mentioned, in the one category it names itself after. Its homepage title is "Shop GhostBed: Luxury Cooling Mattresses & Bedding," and the cooling answer goes to six other brands.

Query, web-search-backed AI answer: "best online mattress brands 2026"

"*WinkBed (the number-one pick), Casper Snow, Leesa, Nectar, DreamCloud, Saatva, Helix...*"

Two of GhostBed's direct competitors, Nectar and Saatva, are named here. GhostBed is not. A third query, "best mattress 2026," returned more than a dozen brands (WinkBed, Helix, Amerisleep, Saatva, Casper, Avocado, Purple, Tempur-Pedic, Sealy, Serta, Beautyrest, Stearns & Foster). GhostBed was in none of them.

GhostBed does surface in one situation: when the query already names it. "GhostBed vs Purple" and "GhostBed vs Saatva" return balanced comparisons, and the engines pull GhostBed's own comparison pages and Tom's Guide. That is brand defense for shoppers who already know GhostBed. It is not discovery for the shoppers who do not. (web-search-backed retrieval, June 22, 2026; sources cited included Sleep Foundation, Mattress Nerd, NapLab, Consumer Reports, RTINGS)

Why this is the most expensive finding in the deck

Category recommendation is the moment of intent, a shopper deciding which brand to buy. GhostBed being absent there is worth more than any single ranking position, because the engine is making the choice on the shopper's behalf and never offers GhostBed as an option. A strong product and a live checkout cannot rescue a recommendation they are never part of. And the gap compounds: each engine that leans harder on live retrieval reads the same listicles and reaches the same competitors.

The fix. Recommendability is downstream of the next four findings. A machine-readable entity, a reputation linked to it, and structured comparison content are what pull a brand into these answers. We also map the specific category sources these engines cite ("best cooling mattress," "best online mattress") and build the placement plan to earn GhostBed into them.

The reputation a machine can read is both weak and unlinked

GhostBed's product pages advertise a 4.8 rating over 10,223 reviews. That number is GhostBed's own first-party review widget. The reputation an independent engine reads off the open web is a different number, and nothing in GhostBed's schema links the two.

```
// On-site, in GhostBed's own product schema (Luxe Hybrid PDP)
AggregateRating: 4.8 from 10,223 reviews (first-party widget)
Organization "GhostBed": sameAs = none

// Off-site, the corpus an engine actually cross-references
Trustpilot: 3.5 / 5 (172 reviews, 82% one-star all time)
BBB (GhostBed): 1.61 / 5 (B+, not accredited, 87 complaints in 3 yrs)
ConsumerAffairs: 1.5 / 5 (120 reviews)
Editorial reviews: 4.1 to 4.7 (Sleepopolis, Mattress Nerd, Mattress Clarity)
Amazon listings: 3.9 to 4.4 (per model)
```

Two separate problems, both fixable

First, the divergence is a credibility risk. As engines increasingly cross-reference Trustpilot, the BBB, and Reddit before they recommend, a self-reported 4.8 sitting next to a 3.5 and a 1.61 reads as unsupported, and the engine discounts it. Second, with sameAs: none, GhostBed is doing nothing to steer which reputation an engine attributes to it, and nothing to separate the consumer-facing GhostBed BBB profile (B+, 1.61) from the parent Nature's Sleep entity (A+). The entity is unmanaged at exactly the layer that decides trust.

The sentiment split matters and works in GhostBed's favor. Editorial reviewers rate the mattress itself 4.1 to 4.7. The negative reviews cluster on the company, returns, restocking fees, the mechanics of the 101-night trial, and warranty responsiveness, not the bed. Costco-channel buyers, who get a frictionless return path, skew positive. The reputation problem is operational and structural, not a verdict on the product. (Trustpilot, BBB, ConsumerAffairs, r/Mattress, observed June 22, 2026)

The fix. What we build and deliver: Organization schema with a populated sameAs that links GhostBed to the profiles it wants attributed and consolidates the BBB entity, plus a review-management spec that targets the unsolicited platforms where the score is lowest. Claiming and responding on those platforms is GhostBed's to execute, and how sentiment moves is ultimately the reviewers'. We own the entity linkage and the playbook; acceptance criteria attach there.

The homepage is strong for humans and blank for machines

The parts of the homepage a person sees are well done. The parts a machine reads to build an entity for "GhostBed" are mostly empty. This is the discoverability tax under finding 1.

```
// ghostbed.com homepage <head> and server HTML, June 22, 2026
<title>Shop GhostBed: Luxury Cooling Mattresses & Bedding | GhostBed</title>    // good
<meta name="description" ...> <meta property="og:title" ...>    present, descriptive
JSON-LD on the homepage:  none    // no Organization, WebSite, SearchAction, or Brand
<h1> in server HTML:  none    // "Our website requires JavaScript to operate properly"
```

Why it matters

The title says "GhostBed," but no structured entity tells a machine what GhostBed is, who makes it, what it is rated, or how its pages relate to each other. The body is JavaScript-rendered, so a crawler that does not execute scripts, which is how much AI retrieval still works, reads the head and a near-empty shell. The good news under this is real: all the major AI crawlers reach the page cleanly. GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot, and Google-Extended each returned a 200. The front door is open; there is just nothing structured inside it.

The contrast, from the same-day competitor recon

Purple's homepage ships a full entity graph: Organization, Brand, WebPage, ItemList, and ContactPoint. An engine assembling "who is GhostBed" has Purple's complete record and GhostBed's empty one to work from.

The fix. Server-rendered Organization, WebSite, SearchAction, and Brand schema on the homepage, with the sameAs links from finding 2, and the page's h1 and key copy in the server HTML so non-JavaScript crawlers can read them. Theme-level and mechanical.

GhostBed wrote the content AI engines quote, then hid it from them

This is the finding that should be encouraging, because the hard part is already done. GhostBed publishes exactly the content these engines pull into comparison and buying answers. It is just unstructured and unmapped, so machines cannot see it for what it is.

```
// ghostbed.com content footprint, from the sitemaps, June 22, 2026
Total /pages:          429
Head-to-head "vs" comparison pages:  36  // ghostbed-vs-saatva, -casper, -helix, -brooklyn-
bedding...
Buying guides (review / best / guide): ~86
Article or FAQPage schema on them:  none
Blog articles in sitemap_blogs:  1
llms.txt: by its own text, "mirrors" agents.md  // a transaction file, not a content map
```

Why it matters

Original, branded comparison and education content is one of the few durable advantages in AI visibility, because it is what engines quote and it is hard to fake. GhostBed has a deep library of it, including its own "GhostBed vs Saatva" and "GhostBed vs Purple" pages, the exact pages an engine wants when a shopper asks for that comparison. But none of it carries Article or FAQPage schema, none of it is mapped in a content-facing llms.txt, and the llms.txt that should do that mapping says, in its own words, that it just mirrors the agent transaction file. The work exists. It is wired to nothing.

The fix. Article and FAQPage schema across the comparison and guide pages, internal links from each to the products it supports, and a rewritten llms.txt that is a genuine content map (the lineup, the cooling technology, the comparisons, the Venus Williams line, trial and warranty terms) kept distinct from the agents.md transaction file.

Product schema is strong; the entity and policy schema around it is missing

This is the dimension GhostBed is closest to winning. Its product pages are genuinely well structured, which is why quotability scores a 70 and not lower. The gap is the layer around the product, not the product itself.

```
// GhostBed Luxe Hybrid PDP, structured data found, June 22, 2026
Product + 7 × Offer (price $1,849 to $3,798, availability InStock)  present
AggregateRating + Review + Rating  present, server-rendered

WebSite / SearchAction:  none
BreadcrumbList (navigation):  none
FAQPage (trial, warranty, shipping, cooling):  none
```

Why it matters

An AI shopping engine can already read GhostBed's prices and variants, which most of the catalog exposes correctly. What it cannot do is answer the questions that come right before a mattress purchase. "What is GhostBed's return policy," "is the trial really 101 nights," "is GhostBed foam or hybrid," and "how does the cooling work" have no structured answer on the page, even though GhostBed has written all of it in prose elsewhere on the site. The product is quotable; the brand and its policies are not.

The fix. Extend the product template that already works with BreadcrumbList and FAQPage (trial, financing, cooling, warranty), and add the WebSite and Organization entity from finding 3. This is incremental on a strong foundation, not a rebuild.

The agent-commerce layer is live, and ahead of the entire category

This is the one dimension where GhostBed is not behind. It is in front. The rails an AI shopping agent needs to find, cart, and check out are live and standards-correct, and none of GhostBed's named competitors run them. We verified the endpoints directly.

```
// ghostbed.com agent surface, June 22, 2026
/.well-known/ucp      live  (UCP profile, version 2026-04-08; checkout, fulfillment, discount)
/api/ucp/mcp          live  (responds over MCP / JSON-RPC; agent-profile handshake required)
agents.md · llms.txt  200   (branded, GhostBed-specific)
sitemap_agentic_discovery.xml  present
Offer pricing for agents  machine-readable · Shop Pay buyer-approved checkout

Saatva · Purple · Nectar:  no ucp, no agents.md, no llms.txt
```

Why it matters, and the honest caveat

GhostBed gets most of this from being on Shopify, so it is a platform advantage rather than a custom build, and the MCP endpoint requires a proper agent-profile handshake before it returns its tools. We are not crediting GhostBed for engineering it. But an agent that arrives reads a live, correct commerce surface, and that is genuinely rare: the three competitors a shopper is most often steered toward instead cannot transact with an agent at all. The rails are the last step of a purchase. Findings 1 through 4 are the reason an agent or a shopper rarely reaches them for GhostBed. Close those, and GhostBed becomes the one brand in its set that an AI engine can both recommend and complete the sale through. That is a strong place to build from.

The lever. Keep the Shopify agent layer current, and make the upstream entity and content legible (findings 1 to 5) so engines and agents choose GhostBed before they ever reach the checkout it has already built.

Competitive position

The same machine-readable signals across the four brands, all checked the same day. GhostBed is the only one that runs a live agent endpoint, and one of two showing an empty homepage entity. It leads the bottom of this table and trails the top.

SIGNAL	GHOSTBED	SAATVA	PURPLE	NECTAR
Homepage entity schema (JSON-LD)	none	Offer only	Org, Brand, WebPage, ItemList	none
Homepage h1 in server HTML	none	1	1	2
Named in AI category answers	no	yes	yes	yes
Live agent endpoint (UCP / MCP)	yes	no	no	no
Branded agents.md + llms.txt	yes	no	no	no

GhostBed and Nectar share the empty-homepage-entity gap, but Nectar still earns the AI recommendation because its category presence and reputation carry it; GhostBed has neither the entity schema nor the recommendation. Purple is the model for the entity layer, a full graph on the homepage. None of the three built the agent rails GhostBed already runs. The reputation comparison sits in finding 2; competitor reputations here are estimates from public signals, consistent with their scorecard columns.

The work

Two phases. Phase 1 is a 30-day stabilization sprint that makes GhostBed machine-readable as an entity, links its reputation, and structures the content it has already written. Phase 2 is ongoing AI visibility work that moves the composite from C toward B over the following months. After Phase 1 we measure against benchmarks set together and decide whether Phase 2 makes sense.

PHASE 1 · 30 DAYS

\$1,000

Stabilization sprint

Six remediations, each closing a finding above.

→ Homepage Organization, WebSite, SearchAction, and Brand schema, server-rendered, with the h1 in server HTML (findings 3, 5)

→ sameAs entity linkage connecting GhostBed to the reputation it wants attributed, plus a review-management spec and BBB-entity consolidation plan (finding 2)

→ Article and FAQPage schema across the 36 comparison and roughly 86 guide pages, with internal links to the products they support (finding 4)

→ llms.txt rewritten as a real content map, kept distinct from the agents.md transaction file (finding 4)

→ BreadcrumbList and FAQPage (trial, warranty, financing, cooling) on the product template (finding 5)

→ Category-source map and placement plan for "best cooling mattress" and "best online mattress," and keeping the live agent layer wired to the new entity (findings 1, 6)

Phase 1 benchmarks, finalized at kickoff.

Composite AI visibility score (target 61 to 72-plus, C to B-range); homepage and product schema-validator coverage including Organization with a populated sameAs; an llms.txt content map live and distinct from agents.md; a measurable lift in the linked external reputation signal; and movement toward presence in AI category answers for cooling and online mattresses. \$1,000 all-in for Phase 1, refunded in full if the benchmarks are not met.

What is ours, what is yours, what the engines decide

We are capability-honest about scope. We build and deliver the schema, the entity linkage, the content maps, the comparison-page structuring, and the placement and review-management specs. Claiming and responding on review platforms, and publishing CMS changes, are GhostBed's to execute. Whether a third-party listicle adds GhostBed, and how sentiment moves, are decided by publishers and reviewers. Acceptance criteria attach only to what we produce.

Phase 2 and beyond

Phase 1 makes GhostBed readable. Phase 2 makes it recommended, and keeps it there. This is the work that is not a one-time fix: continuous AI behavior monitoring per dimension with drift alerts; an entity establishment campaign so "GhostBed" resolves cleanly with its reputation across the engines; getting GhostBed placed in the cooling-mattress and

best-online-mattress sources AI engines cite; and tracking the UCP and agent-commerce changes that move every quarter, where GhostBed already has a head start. Phase 2 scope and pricing are discussed after Phase 1 benchmarks land.

How we built this audit

Every finding came from running specific tools and queries against public surfaces, in an order designed to separate what GhostBed publishes from what machines actually do with it. It is all externally reproducible.

Truth table, frozen first

Before querying any engine, we pulled GhostBed's own machine surface directly: homepage and product <head> and JSON-LD, robots.txt, agents.md, llms.txt, /.well-known/ucp, the live /api/ucp/mcp endpoint, and the product, page, collection, and blog sitemaps. Counts (about 40 products, 429 pages, 36 comparison pages, 1 blog article) come from those sitemaps, pulled June 22, 2026. AI-crawler access was tested by fetching a product page under the GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot, and Google-Extended user agents.

AI behavior and recommendability

We ran category and brand queries through web-search-backed AI retrieval, the same retrieval ChatGPT, Perplexity, and Google AI Mode now use by default, and recorded which brands were named and which sources were cited. GhostBed's absence from the cooling, online, and overall best-mattress answers, and its presence only in queries that already name it, is the basis of finding 1.

Reputation corpus

We pulled GhostBed's third-party reputation from Trustpilot, the BBB, ConsumerAffairs, Amazon, Reddit, and the major mattress-review publishers, and checked whether any of it is linked to the site's entity. The on-site 4.8 is a first-party widget; the off-site corpus (3.5, 1.61, 1.5) and the missing sameAs are finding 2. Amazon per-model figures are live page reads and move over time.

Competitive set

Saatva, Purple, and Nectar were checked the same day on the same signals (homepage entity schema, server h1, AI category presence, agent endpoints, agents.md and llms.txt). Their dimension scores are estimates from those public signals, not full audits, and are labeled as such on the scorecard.

Reproducibility. Every technical finding is verifiable with view-source, the Schema.org validator, and the store's own robots.txt, sitemap.xml, agents.md, llms.txt, /.well-known/ucp, and /products.json. Reputation and recommendability findings are reproducible by running the same public queries. Counts pulled June 22, 2026.

Blackwell Enterprises · Armaan Priyadarshan · Samarjit Deshmukh · Ethan Zhou · Shamit D'Souza · ethanpzhou@berkeley.edu