

Hanz de Fuko

AI Visibility Audit

An external audit of how AI engines and search see Hanz de Fuko today. And what we would do about it.

Prepared for	Hanz de Fuko leadership
Date	June 25, 2026
Scope	Public site, search and AI surface, reputation, agent-commerce readiness
Contact	ethanpzhou@berkeley.edu
Prepared by	Armaan Priyadarshan, Samarjit Deshmukh, Ethan Zhou, Shamit D'Souza
Dartmouth Math & CS · Dartmouth CS · UC Berkeley MET (EECS & Business) · Stanford CS & Biology	

AI visibility scorecard

We graded Hanz de Fuko on five dimensions of how AI engines and search treat the brand. Each dimension is scored 0 to 100 against observed behavior across AI answer engines, the category listicles those engines cite, third-party reputation sources, the brand's own indexed pages, and the agent-commerce surface. Competitor scores are our estimates from the same public signals. Composite is the simple average.

B

76 / 100

HANZ DE FUKO TODAY

Hanz de Fuko is in good shape. AI engines already recommend it by name, the hero product's rating is machine-readable, and the brand runs the full agent-commerce stack. Claymation is named in five of the six engines we tested, and is ChatGPT's outright top pick for "best men's hair clay." The headroom is Discoverability (66): the homepage carries no Organization or Brand entity, so the brand the engines recommend has no machine-readable identity to resolve to, and Hanz holds no claimed reputation profiles. The fixes are entity-level, not a rebuild.

DIMENSION	HANZ DE FUKO	BAXTER	AM. CREW	REUZEL	WHAT IT MEASURES
Discoverability	66	68	70	80	Entity presence, agent discovery files, crawler access, third-party listings
Quotability	74	70	80	70	Schema depth, structured product and policy data, accuracy in machine-readable form
Recommendability	78	76	80	76	Inclusion in AI category recommendations and the listicles engines cite
Transactability	88	80	80	90	UCP and MCP endpoints, payment handlers, agent-readable pricing
Reputation	74	74	78	78	Review aggregator presence and rating, sentiment, links from the entity to that reputation
Composite	76 (B)	74 (C)	78 (B)	79 (B)	Hanz sits mid-pack in a strong set. It wins on recommendation and the agent layer; it trails on the homepage entity, where the brand engines name has no structured identity behind it.

Grading scale: A 90+, B 75 to 89, C 60 to 74, D 50 to 59, F below 50. Competitor scores are estimates from observed public signals, not a full audit of those brands. Reuzel's column is from our full audit of that brand. Method and queries in the methodology section.

The short version

Hanz de Fuko is a well-regarded premium men's styling brand: Quicksand and Claymation are category staples, sold direct and through Amazon, Ulta, and Sephora, on a Shopify store that runs the full agent-commerce stack. Asked for the best men's hair clay, AI engines name Hanz by name, and ChatGPT calls Claymation "the best all-around men's hair clay." The hero product's rating is even server-rendered, which is a step ahead of most of the set. The gap is upstream of all of it: the homepage carries no machine-readable brand entity, so the brand the engines recommend has nothing structured to anchor to, and the reputation lives on channels Hanz has not claimed.

Six findings, ordered by leverage. The last two are strengths.

1. **The homepage has no machine-readable brand entity.** The homepage ships only `WebSite` and `SearchAction` JSON-LD: no `Organization`, no `Brand`. Engines name Hanz de Fuko, but nothing on the homepage tells a machine what the brand is, who makes it, or how its reputation attaches. This is the discoverability gap that caps everything else.
2. **AI engines already recommend Hanz, with one clean gap.** In a six-engine battery, Claymation is named in ChatGPT (top pick), Claude, Gemini, Copilot, and Google's AI Overview. The one miss is Perplexity in incognito, which named Layrite, American Crew, and Suavecito instead.
3. **The product schema works, but duplicates itself.** The Quicksand page serves `Product`, `AggregateOffer`, and `AggregateRating` in the raw HTML, which is good. It also serves two `Product` blocks and two `AggregateRating` blocks on the same page, the kind of duplication that lowers an engine's confidence in which record is canonical.
4. **The reputation is real but unclaimed and unlinked.** 4.3 to 4.6 stars across Amazon, engaged Reddit threads, and a deep YouTube review corpus. But Trustpilot is unclaimed (one review), there is no BBB profile, and none of it is tied to a homepage entity, so engines cannot attribute it confidently.
5. **The rating is already machine-readable, a head start most of the set lacks.** Unlike Reuzel, whose rating hides in JavaScript, Hanz serves its `AggregateRating` in the server HTML. A crawler reading the Quicksand page directly sees the rating. This is the asset to anchor the entity work to.
6. **The agent-commerce layer is live.** A working UCP profile, an MCP endpoint, `agents.md`, and `llms-full.txt`, all returning 200, with every AI crawler reaching the site. An AI agent can find and transact today.

Hanz de Fuko is already recommended and already transactable, with a machine-readable rating most competitors lack. The one thing holding the score down is that the brand the engines name has no structured identity on its own homepage. Give "Hanz de Fuko" a machine-readable entity and link the reputation to it, and the recommendation the engines already make resolves cleanly to the brand's own site.

The homepage has no machine-readable brand entity

AI engines name Hanz de Fuko when a shopper asks for the best hair clay. But when an engine tries to build a record of what "Hanz de Fuko" actually is, who makes it, what it is rated, and how its reputation attaches, the homepage gives it almost nothing. This is the gap that caps the whole scorecard.

```
// hanzdefuko.com homepage JSON-LD, June 25, 2026
Homepage structured data: WebSite + SearchAction // present
Organization entity: none // no name, logo, founder, or sameAs links
Brand entity: none
Title and meta description: present, descriptive · h1 in server HTML: 1
```

Why it matters

The title and the on-page content are fine for a human. What is missing is the structured identity a machine uses to decide that the brand the engines recommend, the products on the PDPs, the 4.x ratings on Amazon, and the YouTube reviews are all the same entity. With no Organization and no sameAs, every reputation signal floats unattached, and the brand is recommended in spite of its own homepage rather than because of it. The good news under this: every AI crawler reaches the site cleanly. GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot, Google-Extended, Amazonbot, and Bingbot each returned a 200. The front door is open; there is just no identity card inside it.

The fix. Server-rendered Organization and Brand schema on the homepage, with sameAs links to the brand's retail, social, and review profiles, so the engines have one entity to resolve every signal to. Theme-level and mechanical, and it is the anchor the rest of the findings attach to.

AI engines already recommend Hanz, with one clean gap

This is where most brands fail and Hanz succeeds. We ran the full six-engine battery, two passes each (from the model's memory with web search off, then with live retrieval on), on "best men's hair clay 2026." Logged-in engines ran in incognito or a temporary chat so saved history could not inflate the result. Claymation is named in five of the six engines, and is ChatGPT's outright top pick.

ENGINE (CLEAN SESSION)	FROM MEMORY (SEARCH OFF)	LIVE WEB RETRIEVAL (SEARCH ON)
ChatGPT (Temporary Chat)	top pick: Claymation	named (\$28, matte, strong hold)
Claude	named	named with Reuzel, Layrite
Gemini	named	named
Copilot (guest)	retrieval only	named with Layrite, Baxter
Google AI Overview	retrieval only	Claymation, top performer
Perplexity (Incognito)	retrieval only	not named

Google AI Overview, live retrieval: "best men's hair clay 2026"

"The best men's hair clay for 2026 delivers texture, volume, and an invisible, reworkable hold. Top performers include Church California Volcanic Clay Pomade, Hanz de Fuko Claymation for all-around versatility, and REUZEL Clay Matte Pomade."



best mer



The best men's hair clay for 2026 delivers texture, volume, and an invisible, reworkable hold without any shine. Top performers for the year include **Church California Volcanic Clay Pomade** for an effortless lived-in look, **Hanz de Fuko Claymation** for all-around versatility, and **REUZEL Clay Matte Pomade** for barbershop-level definition. IH InsideHook +3

Top Hair Clays of 2026

Finding the right hair clay depends on your hair type and desired finish. The top-rated options this year include:

Show more

Videos :




Top 5 Best Hair Clays for Men 2026

YouTube · Review Click

Jan 18, 2026

Discover the **top 5 hair clays for men in 2026**, including Bluman, Seven Potions, CAB's, Lelay, and Smooth Viking.



The best hair clay! Just you wait... a little longer.

YouTube · Robin James

Apr 15, 2026



Best Hair Products for Men 2026 | No More Guessing

Google AI Overview names "Hanz de Fuko Claymotion for all-around versatility" among the year's top performers. ChatGPT, asked from memory, goes further and calls Claymotion the best all-around men's hair clay. Captured June 25, 2026.

The one gap

The single clean miss is Perplexity, run in incognito so no saved history could help or hurt. It named Layrite, American Crew, and Suavecito and left Hanz out. Perplexity leans hard on third-party category guides, so the brand is under-represented in the specific roundups it cites. That is a narrow, addressable gap, not a reputation problem: every other engine, including the from-memory passes, knows Claymotion well.

The fix. Recommendability is already strong; the lever is source authority in the guides Perplexity reads, plus the entity work in finding 1 so the recommendation resolves to a brand the engine can identify and trust.

The product schema works, but duplicates itself

The good news first: the hero product page serves its rating in the raw HTML, which most of the set does not. The issue is that it serves it twice, along with two product records, on the same page. Duplication makes an engine guess which record is canonical, and guessing lowers confidence.

```
// hanzdefuko.com Quicksand PDP, structured data, June 25, 2026
Product      × 2    // two Product blocks on one page
AggregateRating × 2  // two rating blocks on one page
AggregateOffer present // price range, server-rendered
```

Why it matters

A single, clean Product with one AggregateRating is what an engine wants: one unambiguous record to read the price, the rating, and the availability from. Two of each, even with the same values, is the kind of low-grade noise that makes an engine hedge on which to trust, and it can split or dilute how the rating is attributed. This is a deduplication, not a rebuild: the data is right, there is just too much of it.

The fix. Collapse the PDP to a single Product block with one AggregateRating and the AggregateOffer, and tie it to the homepage Organization from finding 1. We deliver the corrected payload; the brand applies it in the theme.

The reputation is real, but unclaimed and unlinked

Hanz has a genuine, engaged reputation. The problem is the same shape as the entity gap in finding 1: the proof exists, but it is not claimed on the channels engines weight, and none of it is tied to a brand entity an engine can resolve.

```
// Hanz de Fuko reputation corpus, live reads, June 25, 2026
Amazon (Quicksand and others): 4.3 to 4.6 / 5
Reddit: active (Claymation how-to threads, "is it still worth?")
YouTube: deep (Quicksand reviews, creator budget breakdowns)
Trustpilot business profile: unclaimed (1 review)
BBB profile: none
```

Why it matters

Amazon and the creator ecosystem rate Hanz well, and the brand is a known name in men's grooming discussion. But Trustpilot, a source several engines weight for trust, sits unclaimed with a single review, there is no BBB profile, and because the homepage has no Organization with sameAs links, an engine cannot confidently connect the strong Amazon and YouTube signal back to the brand. The reputation is real; it is unmanaged at exactly the layer that decides trust.

Reddit shows live, engaged interest: "Hanz De Fuko Claymation, how do you apply this stuff?", "Is Hanz de Fuko still worth?", and "Claymation substitute or better competitor?". This is a brand people actively discuss and compare, which is the raw material the entity and review work turns into a machine-readable signal. (Amazon, Reddit, YouTube, observed June 25, 2026)

The fix. Claim and seed the Trustpilot profile, and link the retail and social reputation through the homepage sameAs from finding 1, so engines attribute the existing strong signal to the brand. Claiming the profile is the brand's to execute; the entity linkage is ours to deliver.

The rating is already machine-readable, a head start most of the set lacks

This is the asset to anchor the entity work to. Where many grooming brands, Reuzel included, lock their rating in a JavaScript widget a crawler cannot read, Hanz serves its `AggregateRating` in the server HTML. A crawler reading the Quicksand page directly sees the rating without running any scripts.

```
// hanzdefuko.com Quicksand PDP, read as ClaudeBot reads it
aggregateRating in raw HTML:  present  // the rating is visible without JavaScript
Contrast, Reuzel PDP raw HTML:  no aggregateRating  // rating hidden in a JS widget
```

Why it matters

The rating is the single fact that most drives a recommendation, and Hanz already exposes it in the form engines read. That is half of what wins the recommendation, and the brand has it. The reason it is filed as a strength with a fix rather than solved outright is finding 3: the rating is served twice, and finding 1: it is not linked to a brand entity. Clean up the duplication and anchor it to an `Organization`, and this becomes a clean, citable, machine-readable rating, which is exactly what pulls a brand into the answers it is not yet named in.

The lever. Keep the server-rendered rating, deduplicate it (finding 3), and link it to the homepage entity (finding 1), so the head start Hanz already has converts into the attribution it is missing.

The agent-commerce layer is live

For most brands this is the hardest dimension. Hanz scores high here because it is on Shopify and the stack is switched on. We verified the endpoints directly.

```
// hanzdefuko.com agent surface, June 25, 2026
/.well-known/ucp      live   (UCP merchant profile)
agents.md · llms-full.txt 200  (agent instructions, full content map)
sitemap_agentic_discovery.xml present
AI crawler access: 200  (GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot, Google-Extended,
Amazonbot, Bingbot)
```

Why it matters, and the honest caveat

An AI agent acting for a shopper can discover the store, read the catalog, and move toward a buyer-approved checkout. The caveat is the same as for any Shopify brand: this is a platform advantage rather than a custom build, and the close competitors run the same stack. There is nothing to fix here. The point is to keep the data the agents read, the brand entity and the deduplicated rating, correct, which findings 1, 3, and 5 handle, so an agent that arrives reads a clean, citable product.

The lever. Keep the Shopify agent layer current, and make the upstream entity and rating legible (findings 1, 3, 5) so engines and agents both resolve cleanly to the brand.

Competitive position

The same machine-readable signals across the four brands, all checked the same day. Hanz is recommended like the rest and exposes its rating, which Reuzel does not. Where it trails is the homepage entity, which none of this set does especially well, and the claimed reputation profiles.

SIGNAL	HANZ DE FUKO	BAXTER	AM. CREW	REUZEL
Named in AI category answers	yes (top in ChatGPT)	yes	yes	yes
aggregateRating in product schema	yes	partial	yes	none
Homepage Organization entity	none	partial	partial	yes
Live agent stack (ucp + agents.md)	yes	yes	yes	yes
Claimed Trustpilot profile	no	n/a	n/a	no

The set is strong across the board: every brand here is named by the engines and runs the agent stack. Hanz’s edge is that its rating is already machine-readable, which Reuzel's is not. Its one clear deficit is the homepage brand entity, where it trails Reuzel. Competitor scores are estimates from the same public signals; the Reuzel column is from our full audit of that brand.

The work

Hanz de Fuko does not need a rebuild. It needs a short entity sprint. Phase 1 is a 30-day pass that gives the brand a machine-readable identity on its homepage, deduplicates the product schema, and links the reputation the brand already earns. Phase 2 is ongoing AI visibility work that holds the position and closes the one engine where Hanz is not yet named. After Phase 1 we measure against benchmarks set together and decide whether Phase 2 makes sense.

PHASE 1 · 30 DAYS

\$1,000

Entity sprint

Four remediations, ranked by leverage, each closing a finding above.

- Server-rendered Organization and Brand schema on the homepage, with sameAs links to the retail, social, and review profiles, so the brand engines name has a machine-readable identity (finding 1)
- Deduplicate the PDP to a single Product with one AggregateRating and the AggregateOffer, tied to the homepage entity (findings 3, 5)
- Link the Amazon, Reddit, and YouTube reputation through the homepage sameAs, and a spec to claim and seed the Trustpilot profile (finding 4)
- Category-source map for "best men's hair clay," targeting the guides Perplexity cites where Hanz is currently absent (finding 2)

Phase 1 benchmarks, finalized at kickoff. Composite AI visibility score (target 76 to 83-plus, holding B and climbing); homepage Organization and Brand schema validator-clean with a populated sameAs; the Quicksand PDP reduced to a single Product and one AggregateRating; and movement toward presence in the Perplexity category answer. \$1,000 all-in for Phase 1, refunded in full if the benchmarks are not met.

What is ours, what is yours, what the engines decide

We produce the exact schema payloads, the entity and sameAs linkage, the PDP deduplication spec, the category-source map, and the re-test that measures movement. Editing the live theme, claiming the Trustpilot profile, and any CMS changes are inside Hanz de Fuko's own systems that the brand executes; we hand over the precise specs and verify the result. Where an outcome depends on a third party, such as a publisher updating a roundup, we flag it and do not promise it.

Phase 2 and beyond

Phase 1 gives the brand a machine-readable identity. Phase 2 holds and extends it: continuous AI behavior monitoring per dimension with drift alerts; an entity establishment campaign so the engines resolve "Hanz de Fuko" cleanly with its rating and reputation; and earning Hanz into the category sources Perplexity cites, the one engine where it is not yet named. Phase 2 scope and pricing are discussed after Phase 1 benchmarks land.

How we built this audit

Every finding came from running specific tools and queries against public surfaces, in an order designed to separate what Hanz de Fuko publishes from what machines actually do with it. It is all externally reproducible.

Truth table, frozen first

Before querying any engine, we pulled Hanz de Fuko's own machine surface directly: homepage and product <head> and JSON-LD, robots.txt, agents.md, llms-full.txt, /.well-known/ucp, the agentic-discovery sitemap, and the product feed (22 products on the first page, 8 collections). AI-crawler access was tested by fetching a product page under the GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot, Google-Extended, Amazonbot, and Bingbot user agents; every one returned a 200. The homepage entity gap in finding 1 and the duplicated product schema in finding 3 came from reading the raw HTML as a crawler reads it.

AI behavior and recommendability

We ran the full six-engine battery (ChatGPT, Perplexity, Gemini, Claude, Google AI Overview, Microsoft Copilot) in headed browser sessions, two passes each: from the model's memory with web search off, then with live retrieval on. Logged-in engines were run in incognito or temporary chat (ChatGPT Temporary Chat, Perplexity Incognito) so saved history did not bias the result; Copilot ran as a guest. Each engine and pass was captured to a screenshot. Claymation was named in ChatGPT (top pick), Claude, Gemini, Copilot, and Google AI Overview; the one miss was Perplexity in incognito, noted in finding 2.

Reputation corpus

We pulled Hanz de Fuko's reputation live from Amazon, Reddit, YouTube, and Trustpilot, and checked whether any of it is linked to a brand entity. The 4.3 to 4.6 Amazon ratings and the active Reddit and YouTube discussion are real; the unclaimed Trustpilot profile, the missing BBB profile, and the absent homepage Organization are findings 1 and 4. Retailer figures are live page reads and move over time.

Competitive set

Baxter of California, American Crew, and Reuzel were checked the same day on the same signals (homepage and product schema, AI category presence, agent endpoints). Their dimension scores are estimates from those public signals, except Reuzel, which is from our full audit of that brand, and are labeled as such on the scorecard.

Reproducibility. Every technical finding is verifiable with view-source, the Schema.org validator, and the store's own robots.txt, sitemap.xml, agents.md, llms-full.txt, /.well-known/ucp, and /products.json. Reputation and recommendability findings are reproducible by running the same public queries. Counts pulled June 25, 2026.

Blackwell Enterprises · Armaan Priyadarshan · Samarjit Deshmukh · Ethan Zhou · Shamit D'Souza · ethanpzhou@berkeley.edu