AI VISIBILITY AUDIT

HUM Nutrition

HUM does the hard part right. Its site is open to the AI assistants, and its product pages carry clean, machine-readable data, so when a shopper names HUM, ChatGPT, Perplexity, Google, and Copilot all read humnutrition.com and get the price and ingredients right. The gap is the question one step earlier. Ask any of them for the best collagen or hair supplement and HUM does not come up. The answer is built from third-party roundups that leave HUM out, and the brand is absent from the moment a shopper is deciding what to buy. This audit shows where that happens and what closes it.

Prepared by Blackwell Enterprises

Team Armaan Priyadarshan (Dartmouth), Samarjit Deshmukh (Dartmouth), Ethan Zhou (UC Berkeley), Shamit D’Souza (Stanford)

Date June 16, 2026

Contact armaan.priyadarshan.29@dartmouth.edu

Summary

HUM Nutrition is in good shape on the parts of AI visibility most brands fail. It lets the AI assistants' crawlers through. GPTBot, the OpenAI and Perplexity search bots, ClaudeBot, and the rest all reach the site and receive the full page. The product pages are rendered on the server and carry well-formed structured data: product, price, in-stock status, return policy, an aggregate rating, and the page FAQs, all in the raw HTML before any JavaScript runs. So a crawler that does not run JavaScript still reads everything it needs.

The effect shows up when you ask the engines directly. Asked the price and ingredients of Collagen Love, ChatGPT, Perplexity, Copilot, and Google all answered $40 with the correct formula and cited HUM's own site. The facts are right and the brand is the source. That is the win, and it is the half of AI visibility that takes engineering to fix when it is broken. Here it is already working.

The loss is recommendability. When the shopper asks by category instead of by name, HUM disappears. Across every engine tested, "best collagen supplement" and "best hair supplement" returned Vital Proteins, Sports Research, California Gold Nutrition, Nutrafol, and others, and never HUM. The engines build those answers from editorial roundups (Rolling Stone, GNC, Health, iHerb) that do not list HUM, so the brand is missing at the exact moment a buyer is choosing.

> **D**
> 66 / 100
>
> Composite across the five dimensions in this audit, measuring the AI channel. The number is held down by one dimension, recommendability, not by the foundation. The schema and crawler access that usually need a rebuild are already working. Closing the gap is corpus and visibility work.

The main points

- The site is readable and the facts are right. On a direct question, every engine tested returned the correct price and ingredients, and most cited humnutrition.com as the source.
- The brand is missing from category answers. "Best collagen supplement" omitted HUM on Perplexity, Google AI Overviews, Google AI Mode, ChatGPT, Copilot, and Claude. Gemini refused the question on safety grounds. The answers that do appear come from roundups that do not include HUM.
- The engines penalize the capsule format. ChatGPT told the user to "choose powder over gummies or capsules," and Google framed capsules as underdosed against powders. HUM's collagen is a capsule, so the frame works against it.
- One crawler is blocked. Bingbot gets a hard 403 at the firewall while the AI crawlers are allowed. The practical harm is smaller than it looks, since Copilot still reached the brand, but it is a closed door that was opened on purpose.
- Nothing is built for shopping agents. No llms.txt, no agent files, no commerce protocol. The product data is clean, so this is a gap to close before agent checkout matters, not a fire.
- Reputation is mixed and the engines reflect it. On-site reviews are strong and deep, Trustpilot sits at 1.8, and the assistants land on "legit, but not essential."

Scorecard

We scored HUM on five dimensions of AI visibility, each from 0 to 100, using the evidence gathered for this audit. The composite is the average.

DIMENSION	SCORE	WHAT IT MEASURES
Discoverability	**75**	Entity presence, agent discovery files, AI crawler access
Quotability	**85**	Machine-readable schema depth, accuracy, structured policies
Recommendability	**47**	Inclusion and framing in AI category recommendations
Transactability	**55**	Agent-commerce protocols, agent-readable pricing, first-party agent checkout
Reputation	**68**	Sentiment in the AI training and retrieval corpus, reviews, safety

How to read the spread

The two high scores and the one low score tell the whole story. Quotability and discoverability are strong because the brand did the technical work: the crawlers are let in and the page data is clean and server-rendered. Recommendability is the outlier at 47, and it is dragging the composite down on its own. That gap is not a site problem. The page reads fine. It is that the third-party corpus the engines lean on for category answers does not include HUM. Transactability and reputation sit in the middle, each for a single clear reason: no agent-commerce surfaces exist yet, and the third-party reputation signals the engines read are weaker than the brand's own reviews.

> By the same five-dimension measure, HUM's standing in classic Google Search would score well into the A range. The site is fully indexed, the product pages return rich results with price, stock, and rating, and the brand ranks for its own terms. The website works. What this audit measures is the newer channel, where being readable is necessary but no longer enough on its own.

Findings, worst first

HIGH 01 The brand is absent from AI category recommendations

This is where the brand loses the most. Asked for the best collagen supplement, every engine that answered returned a list and none of them named HUM. Perplexity returned California Gold Nutrition, Vital Proteins, and Transparent Labs. ChatGPT returned Sports Research, Momentous, and BUBS Naturals. Copilot returned Sports Research, Ancient Nutrition, and Momentous. Claude returned Transparent Labs, Nutricost, and CollagenUP. Google AI Overviews and AI Mode returned Vital Proteins and California Gold Nutrition. Gemini refused the question on safety grounds and recommended nothing. The hair query went the same way, with Nutrafol and Viviscal named and HUM absent. HUM sells a collagen product rated 4.5 across about 200 reviews and a hair product with more than 1,200 reviews, and on a discovery question a shopper never hears either name.

EVIDENCE

Six engines that give recommendations, one query. "Best collagen supplement to buy in 2026" named at least a dozen competitors across the engines and never returned HUM Collagen Love. The answers cited Rolling Stone, Fortune, GNC, Health, and iHerb roundups, none of which list HUM. Gemini, the seventh, declined to recommend supplements at all.

HIGH 02 The engines penalize the capsule format

The category answers do not just omit HUM, they carry a frame that argues against it. ChatGPT's buying rules told the user to "choose powder over gummies or capsules because you can realistically get a full 10 to 20 g serving." Google's overview said a capsule "delivers a much smaller dose of collagen compared to powder-based supplements, which provide much higher yields for a similar price." HUM's Collagen Love is a three-capsule dose, so the engines' default rule of thumb reads it as underdosed before the formula is even considered. The brand has a real answer to this (it pairs collagen with vitamin C, hyaluronic acid, and other actives rather than competing on grams) but that answer is not reaching the model.

CHATGPT, VERBATIM

"A few buying rules: choose powder over gummies or capsules because you can realistically get a full 10 to 20 g serving; choose hydrolyzed collagen peptides."

MEDIUM **03 Bingbot is blocked at the firewall**

HUM's Cloudflare firewall returns a hard 403 to Bingbot on every path, on both the apex and www, while it lets every AI assistant crawler through. There are no AI rules in robots.txt, so this is a deliberate firewall action aimed at Bing specifically. The expected casualty is Microsoft Copilot, which is fed by the Bing index. In testing, the harm was smaller than the block implies: Copilot still answered HUM's price and ingredients correctly and cited the official site, apparently from cache or a path that is not Bingbot. Still, this is a closed door that was opened on purpose, it puts Bing web results and freshness at risk, and reopening it costs nothing.

> EVIDENCE
>
> Requests that differed only in the user-agent returned 200 for a browser and for GPTBot, and a Cloudflare-served 403 for bingbot, repeatably.

MEDIUM **04 Nothing is built for shopping agents**

None of the surfaces an AI shopping agent looks for exist on the site. There is no llms.txt or agents file, nothing under /.well-known (no commerce protocol, no agent card, no plugin manifest), and no MCP endpoint or agent checkout. An agent that wanted to buy HUM would have to go to a retailer, and the sale and the customer go with it. The product data itself is clean and agent-readable, which is the hard part, so this is a gap to close before agent checkout becomes common, not a fire today.

> EVIDENCE
>
> /llms.txt, /agents.md, /.well-known/ucp, /.well-known/agent.json, and /.well-known/ai-plugin.json all return 404.

MEDIUM **05 The third-party reputation the engines read is weaker than the brand's own**

HUM's on-site reviews are strong and deep, with ratings from 4.2 to 4.75 and several products past a thousand reviews. The signals the engines reach for tell a softer story. Trustpilot sits at 1.8 out of 5 across 72 reviews, and the BBB profile is rated A- but not accredited. When asked whether HUM is worth it, Perplexity drew on Trustpilot, the BBB, Medical News Today, and the Chicago Tribune and landed on "legit yes, essential no," praising the testing and clean-label certifications while flagging price and the subscription model. The brand does not control this corpus, and right now it skews more negative than the first-party reviews do.

> EVIDENCE
>
> Trustpilot aggregate rating 1.8 / 5 over 72 reviews on the claimed profile; BBB A-, not accredited; on-site catalog ratings 4.2 to 4.75.

STRENGTH 06 The site is fully readable and accurately quoted

This is the foundation, and it is solid. The AI assistants' crawlers are allowed through the firewall, the product pages are server-rendered, and the structured data is canonical and complete. The result is that on a direct question the engines read the brand's own page and get it right. ChatGPT returned the full Collagen Love ingredient table to the milligram and cited HUM. Perplexity, Copilot, and Google all returned the correct $40 price and the subscribe price, and cited the official site. The brand is the source of its own facts, which is exactly where most brands fall down and HUM does not.

EVIDENCE

Raw product HTML carries a well-formed Product node with Offer (price $40, in stock, return policy), AggregateRating, and FAQPage, using the canonical https schema.org context. Google, ChatGPT, Perplexity, and Copilot each returned the price correctly.

1. Agentic-commerce files and protocols

We checked each file and protocol an AI agent uses to find, read, and buy from a store. The product data an agent needs is present and clean. The discovery and checkout surfaces around it are not built yet.

SURFACE	PURPOSE	STATUS ON THE BRAND'S SITE
AI crawler access	Let assistant bots read the site	**Open. GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, PerplexityBot all 200**
WAF crawler policy	Firewall access control	**403 for Bingbot at the Cloudflare edge**
robots.txt AI rules	Standard crawler opt-out	**None. Only funnel and account paths disallowed**
Product JSON-LD in source	Machine-readable product	**Present, server-rendered, well-formed**
Offer, price, return policy	Agent-readable commerce data	**Present (price, in stock, priceValidUntil, returns)**
OpenGraph product meta	Hints for parsers	**Present in source**
llms.txt, llms-full.txt	Guide written for LLMs	**Absent**
agents.md	Instructions for agents	**Absent**
/.well-known/ucp	Commerce protocol endpoint	**Absent**
/.well-known/agent.json	Agent identity card	**Absent**
/.well-known/ai-plugin.json	Plugin manifest	**Absent**
MCP endpoint	Agent tool interface	**Absent**
ACP (OpenAI, Stripe), AP2 (Google)	Agent checkout and payment mandates	**Not set up**

The split is clean. The top of the table, the data an agent reads, is green. The bottom, the surfaces an agent uses to discover and transact, is missing. Because the product data is already clean, building the missing surfaces is additive work, not a rebuild, and it gets the brand ready before ChatGPT and Perplexity checkout become the default path.

2. How the brand appears in answer engines

We put the same questions to each engine and recorded the answer, the sources, whether the price and ingredient claims were right, and whether the brand appeared in category answers. The results split along one line: HUM does well when named and vanishes when the question is by category.

ENGINE	READS HUM'S SITE	FACTS RIGHT WHEN NAMED	IN CATEGORY ANSWERS	WHERE THE ANSWER COMES FROM
Perplexity	**Yes, cites it**	**Yes**	**No**	brand site, Trustpilot, MNT, roundups
Google AI Overviews	**Yes, cites it**	**Yes**	**No**	brand site, Amazon, Rolling Stone, GNC
Google AI Mode	**Yes**	**Yes**	**No**	Health, Rolling Stone
ChatGPT (search on)	**Yes, cites it**	**Yes**	**No**	brand site, brand pages, PMC, FDA
Microsoft Copilot	**Yes, cites it**	**Yes**	**No**	brand site, Amazon, NIH label database
Claude (search on)	**Facts via retailers**	**Yes**	**No**	Fortune, Rolling Stone, Grove, Amazon
Gemini	**Yes, cites it**	**Yes**	**Refuses, safety filter**	brand site (declines to recommend supplements)
Amazon Rufus	Not captured	Not captured	Not captured	Human-only, Amazon blocks crawlers

The table is uniform in a telling way. Every engine that gives a recommendation reads the brand's own page and gets the facts right, and every one of them leaves the brand out of category recommendations. That rules out the usual suspects. It is not a crawler block, since the page is read. It is not bad data, since the facts are right. It is that the answer to "what should I buy" is assembled from roundups and comparison content where HUM is not present. Two engines are worth a separate note. Claude got the facts right but cited retailers (Grove, Amazon) rather than HUM's own page, even though it is allowed to crawl the site. Gemini refused the category question outright on safety grounds, so for supplements it is not a recommendation surface at all.

HUM against the field

On the collagen question the recurring winners were Vital Proteins, Sports Research, California Gold Nutrition, Ancient Nutrition, and Momentous. Most are powders, and the engines' shared rule of thumb favors powders by gram count, which is the frame that pushes HUM's capsule out of contention. The brand cannot argue its own case in these answers, not because its page is unreadable, but because the page is never the source the engine is summarizing. Getting HUM into that source material is the work of the next phase.

3. What is true, and what confuses the engines

We built a verified table of the brand's products and prices from the rendered pages and the product schema, then checked the engine claims against it. The facts are simple and the engines get them right when they read the page. The confusion is upstream, in how the category is framed.

PRODUCT	PRICE	WHAT IT IS FOR	RATING
Collagen Love	$40	Skin firmness and elasticity, collagen with vitamin C and hyaluronic acid	4.53 / 198
Daily Cleanse	$27	Skin clarity and detox support	4.59 / 2,493
Flatter Me	$27	Digestive enzymes for bloating	4.49 / 2,869
Hair Strong Gummies	$26	Hair strength, biotin and folic acid	4.61 / 1,219
Private Party	$27	Vaginal probiotic	4.72 / 1,892
Gut Instinct	$27	Daily probiotic	4.62 / 1,207
Mighty Night	$40	Overnight skin and sleep support	4.75 / 693
Skin Squad Pre+Probiotic	$40	Clear skin, pre and probiotic	4.71 / 690

Prices and ratings from the brand's own product schema, frozen before any engine was queried. The full catalog of 37 products is in the working file; prices run $12 to $60.

What confuses the engines

- Some products carry a newer name on an older URL. Moody Bird now reads as Hormone Balance, Hair Sweet Hair as Hair Strong Gummies, Glow Sweet Glow as Hyaluronic Glow Gummies. An engine that learned the old name from a third-party source may not connect it to the current page.
- The collagen dose is given in milligrams of a blend, while the products the engines recommend are powders measured in grams. Read without context, the brand looks underdosed, even though it is a different kind of product.
- With the brand's page rarely used as the category source, the engines learn HUM mostly from retailer listings, Trustpilot, and roundups, and some of that is out of date or incomplete.

4. What people say online

This is the material the engines take in about the brand, and it is split. The first-party reviews are strong, the third-party signals are weaker, and the engines weigh the third-party ones because that is what they can reach and trust as independent.

On-site reviews

Deep and positive. Ratings run 4.2 to 4.75 across the catalog, and the most-reviewed products clear a thousand reviews each, with Flatter Me near 2,900 and Daily Cleanse near 2,500. This is a real asset, and it is well-marked up in schema, so the engines see it on a direct query.

Trustpilot, BBB, and editorial

The independent signals are softer. Trustpilot sits at 1.8 out of 5 over 72 reviews, the usual low-volume, negative-skewed pattern for a direct-to-consumer brand, on a profile the brand has claimed. The BBB profile is rated A- but not accredited. Medical News Today and the Chicago Tribune cover the brand on its testing and clean-label credentials, which is favorable. The engines blend these into a balanced verdict.

> When asked whether HUM is worth it, Perplexity answered "legit yes, essential no." It credited the triple-testing, the Clean Label Project certification, and the third-party testing, and flagged the price and the subscription model. That is a fair read of the corpus as it stands, and it is the corpus, not the product, that sets the ceiling.

What to do

The technical foundation is already in place, so the work points at the corpus and the missing surfaces rather than a rebuild. In order of impact:

Get HUM into the category sources the engines read

This is the highest-leverage move and it maps directly to the lowest score. The engines build category answers from roundups and comparison content (Rolling Stone, GNC, Health, iHerb, and the beauty press), and HUM is missing from them. We can produce the target list, the gap analysis of which roundups rank for which queries, and outreach-ready briefs and product data for each. The publishers decide whether to add the brand, so this is influence rather than a switch we flip, but it is the work that moves recommendability.

Answer the capsule-versus-powder frame on the page and in the corpus

The engines default to grams of collagen and read a capsule as underdosed. HUM's case is that it pairs collagen with vitamin C, hyaluronic acid, and other actives for a different goal. Put that comparison in the product page copy and the FAQ schema, in plain terms the model can quote, and seed it into the comparison content. The goal is for the engine to have the brand's own framing available when it reaches for the powder rule of thumb.

Reopen Bingbot at the firewall

Let bingbot through the Cloudflare rule. If the intent was to block training crawlers, that belongs in robots.txt, not a firewall block on the search index. This is a small change the brand's team makes, it protects Bing web results and Copilot freshness, and it costs nothing.

Publish the AI guide files and decide the agent-commerce stance

Add llms.txt and an agents file with the brand's own product and policy information, so the assistants have a clean source the brand wrote. Then decide whether to expose a commerce protocol surface now or let agent sales run through retailers for the moment. Either way, the product data is already agent-readable, so the brand is not shut out when agent checkout arrives.

Strengthen the third-party reputation signals

The on-site reviews are strong; the gap is the independent sources the engines trust. Work the claimed Trustpilot profile to lift the volume and the average, since the engines read it directly, and keep the testing and certification story present in the press the engines cite.

How this gets measured

Every item above is testable, which is the point. The same seven-engine battery that produced this audit, run before and after the work, shows exactly what moved. The schema and crawler fixes are specs we hand the brand's team. The corpus and reputation work is influence we drive, though publishers and reviewers decide what they publish. The before-and-after re-benchmark is the proof, and it is the part that is hard to reproduce without this measurement harness. That is how a focused engagement from here would be scoped and graded.

Method: passive recon and a product truth table were frozen from the brand's own site before any engine was queried. Crawler access was tested with named user-agents. The live engine battery ran in a headed browser across Perplexity, Google AI Overviews, Google AI Mode, ChatGPT, Microsoft Copilot, Claude, and Gemini on June 16, 2026. Amazon Rufus is human-only (Amazon blocks crawlers) and was not run. Scores are graded against the frozen truth table. This is a draft for internal review.