

Good Molecules

Good Molecules blocks AI crawlers from reading its website. Ask ChatGPT or Perplexity about the brand and the answer comes from Ulta, Target, and Reddit. The brand's own site never appears. Google reads the site without any problem. Good Molecules set up the block and can remove it.

Prepared by Blackwell Enterprises

Team Armaan Priyadarshan (Dartmouth), Samarjit Deshmukh (Dartmouth), Ethan Zhou (UC Berkeley), Shamit D'Souza (Stanford)

Date June 1, 2026

Contact armaan.priyadarshan.29@dartmouth.edu

Summary

Good Molecules has set its web application firewall to refuse the crawlers that AI assistants use. ChatGPT's, Claude's, and Perplexity's crawlers get a 403 or a CAPTCHA they cannot solve, so those assistants cannot read the site. The product pages also carry no data in their raw HTML, and the structured data they load afterward is written incorrectly, so even a crawler that got through would find little to read. The search crawlers, Googlebot and Bingbot, are allowed through, so the AI surfaces that run on a search index, Gemini and Google's AI Mode and AI Overviews, still see the brand. The block lands on the standalone assistants.

The effect shows up in what those assistants say. Ask ChatGPT the price of the brand's main serum and it answers \$12, correctly, while citing Ulta and Target. Claude answers the same way and says why, telling the user it could not read the brand's site and used Amazon and other retailers instead. On a category question like the best affordable dark spot serum, Perplexity drops the brand entirely and returns The Ordinary. The facts survive because the retailers carry them. On the standalone assistants, the brand is no longer the source of its own story.

D

61 / 100

Composite score across the five dimensions in this audit. It measures the AI channel. By the same measure, the brand's standing in classic Google Search would score in the 90s. The gap between those two numbers is what this audit is about, and closing it is a configuration and schema problem that does not require a rebuild.

The main points

- The block lands on the standalone assistants. ChatGPT, Claude, Perplexity, and Copilot cannot read the site and answer from retailers and Reddit. The search-fed surfaces, Gemini and Google's AI Mode and AI Overviews, are allowed through and do show the brand.
- The facts are right, but the brand is not the source. Claude says so out loud, telling the user it could not read the site and used retailers instead.
- On a category question, Perplexity drops the brand and returns The Ordinary, even though the brand's hero serum is cheaper and has far more reviews.
- Bing has the brand in its sources and still got the price wrong, reporting the shipping policy as the price.
- The block sits at the firewall, names individual crawlers, and was set on purpose. Classic Google Search is healthy, so the loss is self-imposed and reversible.

Scorecard

We scored Good Molecules on five dimensions of AI visibility, each from 0 to 100, using the evidence gathered for this audit. The composite is the average.

DIMENSION	SCORE	WHAT IT MEASURES
Discoverability	66	Entity presence, agent discovery files, AI crawler access
Quotability	56	Machine-readable schema depth, accuracy, structured policies
Recommendability	52	Inclusion and framing in AI category recommendations
Transactability	56	Agent-commerce protocols, agent-readable pricing, first-party agent checkout
Reputation	74	Sentiment in the AI training and retrieval corpus, reviews, safety

How to read the spread

Four of the five scores sit below passing, and the cause is the same in each. The standalone AI assistants hit the firewall block, and behind it the source HTML has nothing to read, so on ChatGPT, Claude, Perplexity, and Copilot the brand is erased as its own source. The scores are middling rather than low because a lot still works. The search-fed surfaces read the brand fine through Google's index, the schema exists, and the products sell through major retailers. Reputation scores highest at 74, since the positive Reddit and review coverage is the one part of the brand every assistant can reach. Read together, this is a strong brand with one concentrated problem in the standalone assistants.

By the same five-dimension measure, the brand's classic Google Search performance would score in the A range. Around 995 pages are indexed, and product pages return full rich snippets with price, stock, rating, free delivery, and returns. The website works. One channel was closed at the edge.

Findings, worst first

CRITICAL

01 The standalone AI assistants are blocked from reading the site

The site runs behind AWS WAF on CloudFront, and the firewall blocks the crawlers the standalone assistants use. GPTBot, ChatGPT-User, and Bytespider get a flat 403, which matches the user-agent and blocks them from any address. The rest, including OAI-SearchBot, PerplexityBot, ClaudeBot, and Claude-User, get a CAPTCHA they cannot solve. The search crawlers, Googlebot and Bingbot, are allowed through, which is the whole reason the search-fed AI surfaces still see the brand. robots.txt has no AI rules in it, only two SEO crawlers and a search-parameter path, so the block was placed at the firewall on purpose, and the 403 list names the crawlers one at a time.

EVIDENCE Requests that differed only in the user-agent returned 200 for a browser, 403 for GPTBot, and a 405 CAPTCHA for ClaudeBot, carrying the header `x-amzn-waf-action: captcha`.

CRITICAL

02 Those assistants answer about the brand without reading its site

Asked the price and ingredients of the Discoloration Correcting Serum, ChatGPT answered \$12 and cited Ulta and Target. Claude answered the same and cited Amazon and a few review sites. Perplexity drew from Ulta, Reddit, and YouTube. Copilot's shoppable card pointed at Target. In every case the brand's own site was not a source. The facts were right because the retailers carry them. The brand was simply not the one telling its own story. Claude was the most direct about it. Without being asked how it knew, it told the user it could not read the brand's site at all.

CLAUDE, VERBATIM "The Good Molecules site requires JavaScript to load prices, so I couldn't pull them directly, but based on Amazon and third-party retailers, here's what I found."

HIGH

03 The brand drops out of category recommendations on the assistants

On the standalone assistants, a discovery query costs the brand the most. Asked for the best affordable dark spot serum for 2026, Perplexity returned The Ordinary's Alpha Arbutin, Differin, and AXIS-Y, and did not mention Good Molecules. In a head-to-head it appeared but was cast as the gentler, less potent option, with the recommendation leaning to The Ordinary, drawn from Reddit and blogs. Its own \$12 serum, rated 4.3 across more than 7,600 reviews, never entered the comparison on its own terms. The editorial roundups show the same gap, with the brand absent as a source in the lists from NY Mag, Byrdie, and Vogue. The search-fed surfaces did better here, since Gemini included the brand in a list of affordable brands, but on the assistants a discovery shopper never hears the name.

EVIDENCE Perplexity's category answer named three competitors and omitted the brand. Gemini's affordable-brands list included it.

HIGH

04 Product pages carry no data in their source HTML

The page a browser first receives is a 56 KB shell. Its structured-data block is empty, and the markup holds no price, no schema, and no product information. The product name appears once, inside a login link. Everything loads later through JavaScript. Google runs that JavaScript and sees the full page. AI crawlers do not, so even the ones that are not already blocked would find an empty page.

EVIDENCE A headed browser render showed full content after load. The same page fetched as raw HTML had an empty structured-data tag and no price.

MEDIUM

05 The product schema is malformed and loads only after the page does

Once the JavaScript runs, the page does carry complete Product schema, with the \$12 price, in-stock status, the 4.3 rating across 7,624 reviews, and the ingredient list. The markup is written wrong. The type is lowercase product where it should be Product, and the context points to <http://www.schema.org> where it should be https. Google reads it anyway, which is why the rich snippet appears. Stricter parsers and many AI extractors skip markup that is not canonical. The one place the brand publishes product schema, it publishes it slightly wrong.

EVIDENCE The rendered page used "@type": "product" and a non-canonical context. Google's live result parsed it regardless.

MEDIUM**06 Even Bing, which can read the brand, got the price wrong**

Bing sits on the readable side, with goodmolecules.com in its sources, and it still reported the wrong price. Its answer gave the price as \$35 for orders over \$35 and \$5 below, which is the brand's shipping policy mislabeled as the product price. The serum is \$12. The data Bing needed was on the page, but it is buried behind JavaScript and the schema is malformed, so the one brand fact it pulled, it mangled, on a surface that powers Windows and Edge search. The ingredients it gave were correct.

BING, VERBATIM "Good Molecules Discoloration Correcting Serum is priced at \$35 for orders over \$35, and \$5 for orders below \$35." The actual price is \$12.

MEDIUM**07 An AI agent cannot buy from the brand directly, and the surface that could was not reachable**

None of the surfaces an AI shopping agent uses exist here. There is no llms.txt, no agents.md, no UCP or agent card, no MCP endpoint, and no ACP or AP2 setup. The only structured commerce data the brand keeps feeds Google Shopping. An agent cannot complete a purchase on the brand's site, though it can buy the same product from Ulta, Target, Amazon, Walmart, or Beautylish, so the sale and the customer go to a retailer. Amazon Rufus is the one assistant that could actually place that order, and it could not be reached by the test browser, so the checkout question is still open and worth closing.

EVIDENCE The discovery and well-known paths returned absent or blocked. Rufus was not reachable for the shopping test.

STRENGTH**08 Search and the search-fed AI surfaces are healthy**

Classic Google Search ranks the brand near the top of its category, with full rich snippets for price, stock, rating, and shipping. The AI surfaces that ride that index are healthy too. Gemini, Google's AI Mode, and the AI Overviews all show the brand and cite its own site, which corrects an earlier worry that Gemini was cut off. Reputation is good. Reddit reads the brand as a strong budget option, gentler than The Ordinary and suited to sensitive skin, with a smaller group reporting stinging or breakouts. The hero serum holds 4.3 across more than 7,600 reviews, Trustpilot sits around 3 out of 5, and the BBB page shows nothing alarming. Because the search side ranks well and people like the brand, the work ahead is repair on the assistants, and nothing here calls for rebuilding the brand.

1. Agentic-commerce files and protocols

We checked each file and protocol an AI agent uses to find, read, and buy from a store. Almost every surface is missing or blocked. The one that exists feeds Google, and even that is written incorrectly.

SURFACE	PURPOSE	STATUS ON THE BRAND'S SITE
robots.txt AI rules	Standard crawler opt-out	None. Only two SEO crawlers and a search-parameter path.
WAF crawler policy	Firewall access control	403 for GPTBot, ChatGPT-User, Bytespider. CAPTCHA for the AI search crawlers.
llms.txt, llms-full.txt	Guide written for LLMs	Absent
agents.md	Instructions for agents	Absent
/.well-known/ucp	Shopify commerce protocol	Absent. Not a Shopify site.
/.well-known/agent.json	Agent identity card	Absent
/.well-known/ai-plugin.json	Plugin manifest	Absent
MCP endpoint	Agent tool interface	Absent
ACP (OpenAI, Stripe)	Agent checkout	Not set up
AP2 (Google)	Agent payment mandates	Not set up
Product JSON-LD in source	Machine-readable product	Empty. Loads with JavaScript.
Product JSON-LD rendered	Machine-readable product	Present, written wrong
OpenGraph, product meta	Hints for parsers	Absent
Google Merchant feed	Google Shopping ads	Present

The right-hand column is one column wide. Everything built for AI agents is missing or blocked, and the single working surface serves Google. Since the product sits on five major retailers, an agent that wants Good Molecules buys it elsewhere, and the brand keeps none of that sale. This is the gap that grows as ChatGPT and Perplexity add checkout, and competitors whose platforms ship these features get them with no work at all.

2. How the brand appears in answer engines

We put the same questions to each engine, twice where a browsing toggle existed, and recorded the answer, the sources, the price and ingredient claims, and which competitors came up. The results split along the same line as the firewall.

ENGINE	SHOWS BRAND	CITES BRAND'S SITE	FACTS RIGHT	WHERE THE ANSWER COMES FROM
SEARCH-INDEX-FED, CRAWLER ALLOWED THROUGH				
Gemini	Yes	Yes, primary	Yes	brand site, Ulta, Target
Google AI Mode	Yes	Yes	Yes	brand site, Amazon, Ulta, Target
Google AI Overview	Yes	Yes	Yes	brand site, retailers
Bing	Yes	Yes	No, price wrong	brand site, Target, Amazon
STANDALONE ASSISTANTS, CRAWLER BLOCKED				
ChatGPT	Named only	No	Yes	Target, Ulta
Claude	Named only	No	Yes	Amazon, review sites
Perplexity	Where to buy only	No	Yes	Ulta, Reddit, YouTube
Microsoft Copilot	Named only	No	Yes	Target

The table splits cleanly, and the split is about crawler access rather than model quality. The four surfaces that ride a search engine's index reach the brand, because Googlebot and Bingbot are allowed through the firewall. The four standalone assistants are blocked, so they answer from retailers and Reddit and never cite the brand. Bing is the exception that helps no one, since it reaches the brand and still mangled the price. Two surfaces are unmeasured, Grok behind a login and Amazon Rufus behind an interface the test browser could not drive, and Rufus is the one that matters, because it is where an AI could actually buy the product.

Good Molecules against The Ordinary

On the standalone assistants, The Ordinary is the default pick. Perplexity's open category answer leaves Good Molecules out, and in a direct comparison the brand is cast as the gentler, less potent choice with the recommendation tilting to The Ordinary, drawn from Reddit and blogs. The brand cannot argue its own case on these surfaces because its page cannot be read. The search-fed surfaces treat it better. A full scored comparison across the firewall, schema, and engine tests for each competitor is part of the second phase, and the direction is already clear from this pass.

3. What is true, and what confuses the engines

We built a verified table of the brand's products and prices from the rendered pages and the live Google snippets, then checked every engine claim against it. The facts are simple. The confusion comes from where the engines have to read them.

PRODUCT	SIZE	PRICE	KEY ACTIVES	RATING
Discoloration Correcting Serum	30 ml	\$12.00	tranexamic acid 3% as cetyl tranexamate mesylate, niacinamide 4%, pH 5.5	4.3 / 7,624
Discoloration Correcting Serum	75 ml	\$25.00	same formula	4.3 / 7,624
Niacinamide Brightening Toner	120 ml	\$14	niacinamide	4.6 / 3,436
Hyaluronic Acid Serum	30 ml	\$6	hyaluronic acid	4.8 / 2,788
Yerba Mate Wake Up Eye Gel	15 ml	\$6	yerba mate, caffeine	4.0 / 4,393
Brightening & Dark Spots Bar	85 g	\$8	soap-free cleansing bar	4.5 / 883

Shipping is free over \$35, with a \$5 flat rate below that, and returns are free for 60 days. Orders ship FedEx in two days or USPS in three to five. The brand sells on price and price stability, exact ingredient percentages, a gentle approach, and vegan, cruelty-free formulation.

What the brand cares about

The brand leads with affordability and steady prices, ingredient transparency down to the percentage, a gentle but effective approach, and a vegan, cruelty-free formulation. Most people know it by the Discoloration Correcting Serum.

What confuses the engines

- The same serum comes in two sizes, \$12 for 30 ml and \$25 for 75 ml, and Target lists it at \$11.99. An engine that reads one listing can attach the wrong size to a price, and the number drifts between retailers.
- The 3% tranexamic acid is listed by its chemical form, cetyl tranexamate mesylate. One source calls it tranexamic acid and another calls it the ester, and a reader can take them for two ingredients. The brand's own page would settle this, and it cannot be read.
- With the brand page blocked, the engines learn the product from retailer ingredient lists, Reddit, and review blogs, none of which is the brand speaking, and some of which is out of date.

- The Reddit view that the brand is weaker than The Ordinary is being repeated by Perplexity as close to a verdict. The brand has its own efficacy data, but with its page blocked it has no way to put that in front of the model.
- Where the structured data does exist, the markup errors send strict and AI parsers back to third-party text.

4. What people say online

This is the material AI takes in about the brand, and it is the one area in good shape. The brand just does not control any of it.

Reddit

On Reddit the brand reads as a strong budget option, with formulas people find more elegant and gentler than The Ordinary and well suited to sensitive skin. A smaller group reports stinging, redness, or breakouts. The engines repeat this framing almost word for word, so Reddit sets the tone of what AI says about the product.

Trustpilot, BBB, and safety

Trustpilot sits around 3 out of 5, with a layer of review blogs around it. The BBB page shows nothing alarming, which matters because an earlier audit found a BBB F surfacing inside ChatGPT, and this brand carries no such liability. Talk of PFAS and clean formulation is mild, and a documented denial ranks alongside it.

YouTube and editorial lists

YouTube reviews are plentiful and mostly positive, and the engines increasingly cite them. The editorial roundups are the weak spot. The brand does not appear as a source in the general category lists, the same absence that shows up in the AI recommendations.

The brand's reputation is good, and AI is passing it along faithfully. It is also the only part of the brand AI can reach right now, since the brand's own pages are blocked. The reputation is an asset the brand does not control, and it cannot be steered until those pages can be read.

What to do

Reopen the AI crawlers at the firewall. Let the retrieval and grounding crawlers through, namely OAI-SearchBot, ChatGPT-User, PerplexityBot, Perplexity-User, Claude-User, Claude-SearchBot, and Google-Extended. If the goal is to keep training crawlers out, do that in robots.txt rather than with a firewall block. This is the change that matters most, and it undoes the absence on its own. Check first whether the block was intended. The hand-built deny list suggests it was, and if so the brand is giving up live AI visibility for a training-data stance it could take far more cheaply.

Render the product data on the server. Put the structured data, the price, the stock status, and the shipping and returns terms into the HTML the server sends, so a crawler that does not run JavaScript still sees them. Today that data only appears after the page loads.

Fix the schema. Use the correct Product and AggregateRating types and the https schema.org context, and add OpenGraph and product tags. This lets strict and AI parsers read what Google already reads.

Publish the AI guide files. Add llms.txt, llms-full.txt, and agents.md with the brand's own product and policy information, so assistants have a clean source the brand wrote.

Decide the agent-commerce stance. On a custom platform the agent-commerce protocols do not come for free. Choose whether to adopt ACP, AP2, or an MCP surface, or to let the agent sale run through retailers. Either way, make the product data agent-readable now, so the brand is not shut out when these features arrive.

Turn the Reddit goodwill into recommendations. Once the page can be read, a \$12 serum rated 4.3 across more than 7,600 reviews is strong enough to earn a place in category answers. Pair the firewall fix with a Wikipedia article and a clean Wikidata entry, and with the work to get the brand into the editorial roundups it is missing from.

Put the brand's efficacy story on the page. Once the page can be read, keep the clinical results and the gentle but effective positioning in the page and in the schema, so the model can use the brand's account alongside Reddit's.

Engagement

The team

Armaan Priyadarshan (Dartmouth, Mathematics and Computer Science), **Samarjit Deshmukh** (Dartmouth, Computer Science), **Ethan Zhou** (UC Berkeley M.E.T., EECS and Business), and **Shamit D'Souza** (Stanford, Computer Science and Biology). Blackwell Enterprises is a four-person consulting team.

How the engagement works

The work runs in two phases, and Good Molecules decides on the second after seeing the results of the first.

Phase one reopens the AI channel. It covers the firewall change, the server-side schema, the schema corrections, and the guide files, with benchmarks agreed at the start and measured at delivery. The benchmarks are crawler reachability, schema validation, and how often the brand is cited across the five engines. The sprint starts within five business days of kickoff.

Phase two builds the presence back. It covers ongoing measurement across the five engines, a Wikipedia and Wikidata entry, the work to recover the category recommendations, and the agent-commerce decision. It is scoped once the phase one results are in.

Price

\$1,000

The fee for phase one, the work above. Good Molecules pays the full \$1,000 upfront, refunded in full if the work does not meet the benchmarks agreed at the start. The fee covers two rounds of revisions. Any third-party firewall, CDN, or hosting charges on the brand's own systems sit outside this and are billed at cost or paid directly. Phase two is priced once the phase one results are in.

Everything in this audit can be reproduced with public tools. curl, a browser, the Schema.org validator, Google Search, and ordinary ChatGPT and Perplexity accounts. We used nothing beyond customer access. The crawler tests, response headers, rendered page, engine transcripts, and reputation captures are available on request.

Contact. armaan.priyadarshan.29@dartmouth.edu