

Good Molecules

Since the audit, you fixed the two things that were breaking AI visibility at the root. The firewall now lets the AI assistants reach the site, and your product pages carry real structured data in the server HTML. We confirmed both this morning. That clears the foundation, so the work now is enrichment, not a rebuild. This brief shows exactly what your structured data looks like today, the gaps that are left, and what we do in week one.

Prepared by	Blackwell Enterprises
Team	Armaan Priyadarshan (Dartmouth), Samarjit Deshmukh (Dartmouth), Ethan Zhou (UC Berkeley), Shamit D'Souza (Stanford)
Date	June 17, 2026
Contact	armaan.priyadarshan.29@dartmouth.edu

What changed since the audit, confirmed

You noted two changes on your end. We re-ran passive recon this morning, curling the production site with each crawler’s published user-agent and reading the raw server response. Both changes are live.

1. The firewall now lets the AI assistants in

This was the headline finding of the audit. At audit, the standalone assistants were blocked at the firewall and answered from Amazon, Ulta, Target, and Reddit instead of your site. The crawlers that power ChatGPT, Claude, and Perplexity now return HTTP 200.

CRAWLER (USER-AGENT)	AT AUDIT	JUNE 17
GPTBot (OpenAI index)	Blocked	200 allowed
OAI-SearchBot	Blocked	200 allowed
ChatGPT-User (live fetch)	Blocked	200 allowed
ClaudeBot	Blocked	200 allowed
Claude-User (live fetch)	Blocked	200 allowed
PerplexityBot	Blocked	200 allowed
Perplexity-User (live fetch)	Blocked	405 still blocked

Two follow-ups for your side. Perplexity-User, the agent Perplexity uses for live fetches, is still blocked, so that one user-agent needs the same allowlist entry. And spoofed Googlebot and Bingbot strings return 405, which is likely the firewall correctly rejecting unverified bots, but since the Google and Bing surfaces were your one working channel at audit, we should confirm the verified search crawlers still pass.

2. Product pages now carry server-side structured data

Also confirmed. Product pages serve JSON-LD in the raw HTML with no JavaScript needed: Product, Offer with price and in-stock status, and an aggregate rating. That is the direct fix for the audit finding where Claude could not pull your prices and fell back to Amazon. You also shipped a strong llms.txt, which was effectively absent at audit. Good base to build on.

Your structured data today: full evaluation

We read the JSON-LD on ten product pages across categories, plus the homepage and a category page. The data is uniform across the catalog, so every fix below is one template change that propagates to every product, not page-by-page work.

LIVE AND CORRECT

Product (name with size, image, rich description, brand), Offer (price, priceCurrency USD, availability in-stock, item condition new), AggregateRating with real volume (4.0 to 4.8 across the sample, up to 7,627 reviews), and an OnlineStore node on every page. It is real, server-rendered, and consistent.

Gaps we will close, by impact

PRIORITY 1 No product identifiers (sku, gtin, mpn)

The Product node has no GTIN, UPC, or SKU. GTIN is the key an engine uses to know your page and the Amazon, Ulta, or Target listing are the same product. Without it, engines cannot confidently match your page to the retailer listing, which is why they fell back to retailer prices in the audit. Adding GTINs is what lets your own page win the price.

PRIORITY 2 Size and variant relationships are not modeled

Each size is its own page with a single Offer. There is no ProductGroup tying the sizes together, so an engine on the 30ml page has no signal that other sizes exist or what they cost. This is the residual of the 12ml versus 30ml price confusion.

PRIORITY 2 Offer is missing commerce fields you already have

The Offer has no shipping details, return policy, price-valid-until, or offer URL. Your llms.txt already states the shipping tiers and you have a return policy. Putting them in OfferShippingDetails and MerchantReturnPolicy completes the offer for agentic checkout and Merchant-style listings.

PRIORITY 3 Aggregate ratings only, no review text

You expose the aggregate rating but no individual Review objects, though you have thousands of reviews. Engines quote review text in answers, so surfacing a sample feeds how they describe the product.

PRIORITY 3 Thin homepage entity schema

The homepage carries only an OnlineStore node. There is no sameAs linking your social and retailer profiles, no contactPoint, and no WebSite search action. sameAs is how an engine grounds “Good Molecules” as one confident entity. Low effort, high value.

PRIORITY 3 No breadcrumbs, no collection lists, thin social meta, missing files

No BreadcrumbList on product pages, no ItemList on category pages, and only og:site_name for social meta. Separately, sitemap.xml returns 404, llms-full.txt, agents.md, and the /.well-known/ endpoints are still absent, and there is no FAQPage markup or on-page FAQ.

Worked example: niacinamide serum

Built from your live data. Values marked “you supply” are data only you have.

Live today

```
{
  "@context": "https://schema.org",
  "@type": "Product",
  "name": "Good Molecules Niacinamide Serum 30ml",
  "brand": { "@type": "Brand", "name": "Good Molecules" },
  "image": "https://.../prod_p.jpg",
  "description": "<h5>About...</h5> ...",
  "offers": {
    "@type": "Offer",
    "priceCurrency": "USD",
    "price": 6.0,
    "itemCondition": "https://schema.org/NewCondition",
    "availability": "https://schema.org/InStock"
  },
  "aggregateRating": { "@type": "AggregateRating", "ratingValue": 4.6, "reviewCount": 1751 }
}
```

A clean minimum-viable Product. Everything below is additive to this same block.

Proposed

```
{
  "@context": "https://schema.org",
  "@type": "Product",
  "name": "Good Molecules Niacinamide Serum",
  "sku": "GM-NIA-30", // you supply
  "gtin13": "0860000000000", // you supply (UPC / EAN)
  "brand": { "@type": "Brand", "name": "Good Molecules" },
  "image": "https://.../prod_p.jpg",
  "description": "10% niacinamide serum that refines tone and texture...",
  "url": "https://www.goodmolecules.com/s/good-molecules-niacinamide-serum-30ml",
  "category": "Skincare > Serums",
  "offers": {
    "@type": "Offer",
    "url": "https://www.goodmolecules.com/s/good-molecules-niacinamide-serum-30ml",
    "priceCurrency": "USD",
    "price": 6.0,
    "priceValidUntil": "2026-12-31",
    "itemCondition": "https://schema.org/NewCondition",
    "availability": "https://schema.org/InStock",
    "shippingDetails": {
      "@type": "OfferShippingDetails",
      "shippingRate": { "@type": "MonetaryAmount", "value": 0, "currency": "USD" },
      "shippingDestination": { "@type": "DefinedRegion", "addressCountry": "US" },
      "deliveryTime": { "@type": "ShippingDeliveryTime",
        "transitTime": { "@type": "QuantitativeValue", "minValue": 3, "maxValue": 5, "unitCode":
"DAY" } }
    },
    "hasMerchantReturnPolicy": {
      "@type": "MerchantReturnPolicy",
      "applicableCountry": "US",
      "returnPolicyCategory": "https://schema.org/MerchantReturnFiniteReturnWindow",
      "merchantReturnDays": 30 // confirm your window
    }
  },
  "aggregateRating": { "@type": "AggregateRating", "ratingValue": 4.6, "reviewCount": 1751 },
  "review": [ { "@type": "Review", "reviewRating": { "@type": "Rating", "ratingValue": 5 },
    "author": { "@type": "Person", "name": "from your review feed" },
    "reviewBody": "pull a sample from your on-site reviews" } ]
}
```

Plus a ProductGroup to tie the sizes together, a BreadcrumbList on the page, and a homepage Organization node with sameAs. Same pattern, applied catalog-wide.

Week-one plan

What we deliver this week

Drafts ready to deploy on your platform.

- Corrected and enriched Product and Offer JSON-LD template, with the niacinamide serum worked end to end as the reference, plus the ProductGroup variant model
- Homepage Organization and WebSite schema with sameAs, contactPoint, and a search action
- OfferShippingDetails and MerchantReturnPolicy blocks built from your published terms
- BreadcrumbList for product pages and ItemList for category pages
- OpenGraph and product meta tag spec
- llms-full.txt and agents.md drafts, building on your live llms.txt
- The first FAQ set keyed to real skincare query patterns, with FAQPage markup

What we need from you

- GTIN/UPC and SKU per product from your product system (unblocks Priority 1)
- Your exact return window and policy text (completes Priority 2)
- Access to or a sample export of your review feed (enables the review markup)
- Platform owners to deploy the schema and generate sitemap.xml

How we measure it

We re-check the structured data and crawler access after each deploy and validate it against the schema.org and rich-results requirements, so the before and after is concrete at the technical layer. Live engine answer testing is a separate step we schedule with you, not part of this week.

SCOPE, SO EXPECTATIONS ARE CLEAR

We draft and hand over every artifact above. Deploying to your server and supplying GTINs and return text is yours to execute. How third-party listicles and retailers respond is decided by them; we guide that work but do not promise their behavior.

Method: passive recon and structured-data read on June 17, 2026. Curl with each crawler's published user-agent against www.goodmolecules.com, following redirects; HTTP status and JSON-LD taken from raw server responses across ten product pages, the homepage, and a category page. No live engine testing was run for this brief.