

BLACKWELL ENTERPRISES

Public Goods AI Visibility Audit

An external audit of how six AI platforms see Public Goods today. And what we would do about it.

Prepared for Public Goods leadership

Date May 27, 2026

Contact armaan.priyadarshan.29@dartmouth.edu

Prepared by Armaan Priyadarshan, Samarjit Deshmukh, Ethan Zhou, Shamit D'Souza

Dartmouth Math & CS · Dartmouth CS · UC Berkeley MET (EECS & Business) · Stanford CS & Biology

AI Visibility Scorecard

We graded Public Goods on five dimensions of how AI engines treat the brand. Each dimension scored 0 to 100 against observed AI behavior across ChatGPT, Reddit, Trustpilot, BBB, and the brand's own indexed pages. Composite is the simple average.



Public Goods is the lowest-scoring brand in its competitive set, trailing the average by 27 points. The competitive average across Grove (B+), Thrive (B-), and Blueland (B+) is 79. Public Goods is at 52, the only sub-B grade in the category.

The transactional infrastructure is genuinely strong. Three of the five dimensions that drive AI behavior toward customers (quotability, recommendability, reputation) are failing. The brand pays the cost on every AI shopping query in the category.

Dimension	Public Goods	Grove	Thrive	Blueland	What it measures
Discoverability	75	85	90	75	Entity presence, agent discovery files, crawler access
Quotability	40	75	90	80	Schema depth, structured policies, accuracy in machine-readable form
Recommendability	30	85	90	90	Inclusion in AI category recommendations, credentials, listicle presence
Transactability	80	80	30	75	UCP, MCP, payment handlers, agent-readable pricing
Reputation	35	80	65	80	Sentiment in AI training corpus, BBB, review aggregator presence
Composite	52 (D)	81 (B+)	73 (B-)	80 (B+)	

Grading scale: A 90+, B 75-89, C 60-74, D 50-59, F below 50. Methodology and per-prompt evidence available in the appendix to this deck.

The short version

Public Goods built sophisticated infrastructure for AI agents to transact with the store. AI engines never get that far. They do not recommend the brand in category queries. When asked about it directly, they surface a fraud-adjacent reputation, a deprecated offer the brand discontinued last year, and a shipping policy that contradicts itself across four pages.

Six findings, ordered worst first.

1. **BBB F rating with nine unanswered complaints surfaces in ChatGPT responses**, alongside corroborating "scam" and "fraud" language from Reddit and Trustpilot. Three independent platforms carry the same complaint with the same wording: unauthorized membership renewals years after a single purchase. ChatGPT cites PublicGoods.com by domain when surfacing the BBB record.
2. **Public Goods is absent from every AI-synthesized "best sustainable brands" ranking we ran**. ChatGPT with web search names Blueland, Grove, Seventh Generation, and Dropps as category leaders. Public Goods is not mentioned. Blueland is a smaller company than Public Goods.
3. **A deprecated 40% membership offer is actively indexed and surfaces in ChatGPT responses** to prospective new members. The page calls 40% the "best available discount." That offer was discontinued in April 2025. AI engines are bait-and-switching prospects with the brand's own indexed content.
4. **The free shipping policy yields six different correct answers** across three customer cohorts and two contradicting source pages. ChatGPT noticed and flagged the internal contradiction unprompted in one test.
5. **Member pricing is structurally invisible across every machine-readable surface**. The membership's actual value proposition cannot be quoted by any AI agent that does not render JavaScript. Public Goods loses every head-to-head comparison this way. ChatGPT frames the brand as the pricier option versus Grove and Thrive because the discount cannot be quantified.
6. **PDP schema is materially thinner than Grove and Blueland**. No SKU. No per-variant offers. No FAQPage schema on the 31 accordion FAQs already on the page.

ChatGPT, Perplexity, and Google AI Mode all moved to live web retrieval as default behavior in 2026. Every additional query that hits live retrieval rather than training data widens the gap between Public Goods and the competitive set. The damage compounds weekly.

AI engines tell prospects Public Goods is a scam

The image shows two side-by-side screenshots of the PublicGoods.com website. The left screenshot is the 'BUSINESS PROFILE' page, which includes contact information for PublicGoods.com (Headquarters: PO Box 838, New York, NY 10008-0838), a BBB Accreditation & Rating section stating 'PublicGoods.com is NOT a BBB Accredited Business', and a BBB Rating of 'F'. The right screenshot is the 'BBB Accreditation & Rating' page, which provides more details about the BBB process, a BBB Rating of 'F', and a section for 'Latest Reviews' featuring three negative reviews from Matthew R, Linnop C, and Jordan R, all mentioning issues with annual subscription charges and lack of response to complaints.

F rating. NOT BBB Accredited. 9 unanswered complaints. 4 unresolved. File opened 5/29/2020. Reviewer language: "DO NOT TRUST", "complete scam", "fraudulent annual subscription."

Three independent platforms, same complaint, same language

"Wife purchased items 3-4 years ago. We have been charged \$65 annually for an 'annual subscription' service." (BBB)

"Fraudulent membership fees charged years after one purchase. In 2025 I was charged \$79 for a membership I never authorized." (Trustpilot, 1-star)

"my bank caught it and flagged it as fraud." (Reddit r/ZeroWaste, in a thread titled "PSA: Public Goods is a scam")

Additional Trustpilot 1-star titles, verbatim: *"Fraudulent Cancellation Process"* (canceled in January, charged March 18). *"YOU CAN NEVER CANCEL!! Surreptitiously opted me into a membership." "Absolutely HORRENDOUS customer service... they gaslight you... I believe the 5 star reviews are not real."* The word "Fraudulent" appears in two separate review titles. Reddit r/Scams also carries a 2024 thread: "Public Goods charged my card 4 years after purchase."

What ChatGPT tells a prospect today

Prompt with web search: *"Public Goods customer complaints."* Sources cited: bbb.org, trustpilot.com, reddit.com, publicgoods.com.

"BBB shows PublicGoods.com with an F rating, citing failure to respond to 9 complaints. Multiple 2025-2026 complaints about customers claiming they did not knowingly sign up for an annual membership."

Public Goods does not exist in AI category awareness

Prompt: "Best sustainable household brands." ChatGPT with web search.

ChatGPT named four brands as category leaders, each with citable credentials:

- Blueland (best overall, citing B-Corp certification)
- Grove Collaborative (best one-stop, citing range and certifications)
- Seventh Generation
- Dropps

Public Goods was not in the main list. Not in honorable mentions. Not in a "consider also" section. The brand does not appear at any depth of the response.

Why this is the most expensive finding in the deck

Public Goods appears in 8 of 40 human-written third-party listicles we audited. The brand is in human consideration sets. It is not in AI consideration sets. As AI shopping recommendations replace human listicles in buyer journeys, this gap closes the funnel before it ever opens.

The brands ChatGPT named all have something to cite. Blueland and Grove have B-Corp certification. Public Goods has no surfaced credential. Grove (Q111170467) and Thrive (Q27150214) have Wikidata entities. Public Goods has none. The brand name also collides with the economics concept "public good," which dominates the SERP: four of the top eight organic Google results for "Public Goods" are the econ concept (Wikipedia "Public good", econlib, Investopedia, Khan Academy).

Three further signals AI engines weight

Merchant-only Knowledge Panel. Google's Knowledge Panel for Public Goods exists, but it's a commerce/merchant card (free delivery threshold, returns policy, payment options). Zero entity facts: no founder (Morgan Hirsh, founded 2016), no HQ, no founding date, no CEO. AI engines pulling "what is Public Goods" from the Knowledge Graph get nothing about the brand.

Negative Reddit ranks page 1 of Google for the brand's own name. Search "publicgoods.com": position 7 is the r/ZeroWaste "PSA: Public Goods is a scam" thread. Search "Public Goods brand": position 3 is the r/BuyItForLife "Is Public Goods BIFL?" thread (upvote ratio 0.27). AI engines that crawl SERP for citations pick these up.

Absent from head-to-head comparison content

The top-ranking "best sustainable brand vs" article on Google is titled "Grove vs Blueland vs CleanCult vs Thrive vs Honest." Five brands compared in the category, none of them Public Goods. This is what ChatGPT pulls from when synthesizing comparison answers. Closing the gap is editorial and entity work, not Shopify theme work.

AI is bait-and-switching the brand's own prospects

The page at `/pages/legacy-membership` contains this sentence in its static, server-rendered HTML:

"automatically enrolled in our exclusive VIP Member program. This gives you the best available discount of 40% off most categories (30% off food items, regular pricing on frozen fish and plants), a special benefit..."

"Best available discount" implies the 40% rate is currently the best on offer. It is not. The April 2025 restructure moved new members to a flat 20%. Legacy members were grandfathered into the 40%. Nobody can sign up for that rate today.

Three things compound the problem:

- The page has no `noindex` meta directive or X-Robots-Tag header.
- The canonical URL is self-referencing. Public Goods is declaring the page the canonical version to index.
- It is present in `sitemap_pages_1.xml`. The site actively submits it to Google.

What ChatGPT does with it

Prompt with web search: *"What discount do Public Goods members get?"*

"20% off for new members. One caveat: Public Goods says legacy VIP members may keep a higher discount of 40% off most categories, 30% off food, and regular pricing on frozen fish and plants."

Source cited: publicgoods.com

ChatGPT presents the deprecated offer as a live "caveat" to a new prospect. The prospect signs up, gets the 20%, and learns the 40% rate was never available to them. The brand's own indexed content is producing the disappointment.

WHY THIS GETS WORSE OVER TIME

Every additional AI engine that adds live web retrieval is another engine pulling this page. The page is more useful to AI engines than to humans. Traffic to it is increasingly bots, which means its downstream impact is increasingly on AI responses, not on real customers seeing it directly.

The membership's actual value is invisible to AI

The entire reason to buy the membership is the lower price. That number exists in no machine-readable field anywhere on the site.

Three surfaces, all empty

JSON-LD on the natural shampoo PDP. One flat `offer` at \$19.95. No `priceSpecification`. No member tier. No eligibility logic. No SKU. No GTIN.

The `/products/natural-shampoo.json` endpoint. Rich data: real SKUs, real GTINs, inventory counts, full variant arrays, marketing copy. Zero member or tier pricing fields. The endpoint exposes a "Member Discount" tag as a string but no numeric discount.

Static HTML on the PDP. Twelve `<span class="member-price">` elements. All empty. The actual member price is computed and injected by JavaScript at render time.

The competitive consequence we measured

Prompt: *"Public Goods vs Grove Collaborative."* ChatGPT with web search.

ChatGPT frames Public Goods as the more expensive membership: \$65/year against Grove's \$29.99/year. The 20% discount that would offset this cannot be quantified because the member price is invisible. Public Goods loses every head-to-head this way. Same dynamic with Thrive (\$59.95/year).

THIS IS NOT A GENERIC SHOPIFY THEME ISSUE

Member pricing as eligible-buyer schema requires custom `UnitPriceSpecification` with `eligibleCustomerType`, integrated with the Shopify customer-tag system, populated in both JSON-LD and the product JSON, and rendered server-side rather than client-side. No Shopify app does this out of the box for membership stores. The pattern requires specific schema knowledge and bespoke implementation work.

The shipping policy contradicts itself across four pages

Three customer cohorts. Two policy dimensions. Two source pages that disagree. Six possible correct answers, all crawlable, none mappable to the prospect.

Customer cohort	Free shipping rule	Source page
Members enrolled after 4/1/2025	Free on all orders, no minimum	Membership info
Members enrolled before 4/1/2025	Free on orders of \$45+	Legacy membership
Non-members	\$45 minimum	FAQ page
Non-members	\$65 minimum	Support page

ChatGPT flags the contradiction unprompted

Prompt with web search: *"Does Public Goods offer free shipping?"*

"Their current membership page also says members get free U.S. shipping on every order with no minimum required. The shipping FAQ on a different page lists the \$45 threshold as universal for non-members, but the support page lists \$65. The current threshold depends on which page is treated as authoritative."

The policy is internally inconsistent enough that an AI engine flags it as a contradiction to the user. The user does not have to notice. ChatGPT volunteers the contradiction as part of the response.

Remediation requires more than picking one page and editing it. The policy must be unified across the FAQ, support, policy page, cart drawer, schema `OfferShippingDetails`, and the contradicting pages must be 301-redirected because AI engines have already cached the old versions.

PDP schema is thinner than competitors

One hero PDP per brand. All three are Shopify storefronts in the sustainability category.

Schema field	Public Goods	Grove	Blueland
SKU in JSON-LD Product	null	real SKU	real SKU
GTIN in JSON-LD	absent	gtin12	absent
Offers granularity	1 flat offer	5 per-variant	2 offers
AggregateRating	yes	no	yes
FAQPage schema	no	no	yes
Review schema	no	no	no
additionalProperty	no	no	yes

Blueland is a smaller company. Their PDP schema is deeper.

The variant offer, concretely

Public Goods natural shampoo has four real variants priced at \$9.95, \$19.95, \$37.95, and \$49.95. The JSON-LD collapses this into one `offer` at \$19.95. An agent comparing this product to Grove sees Grove as a fully itemized 5-SKU catalog entry and Public Goods as one undifferentiated item.

Grove emits all five variants as separate offers in JSON-LD even on a product where every variant is the same price (\$9.99 across five scents). The signal to AI is structural depth, not pricing range. Public Goods's choice not to do this is the kind of decision an engineer makes once and never revisits, costing the brand on every AI parse forever after.

The trajectory: AI behavior is shifting against the brand

We ran every prompt twice on ChatGPT Plus. Pass 1 with web search disabled. Pass 2 with web search enabled. Public Goods's AI visibility is worse in Pass 2 on every axis we measured.

Question	Pass 1 (training data only)	Pass 2 (live web retrieval)
What discount do members get?	"20%, not 40%" (clean)	"20%, but legacy VIP keeps 40% off most, 30% off food, regular pricing on frozen fish" (cites publicgoods.com)
Free shipping minimum?	Flat "\$45"	Three-cohort matrix plus internal contradiction flagged unprompted
Is Public Goods legit?	"Not a scam, but..." (hedged)	BBB F rating, \$79 unauthorized charges, "scam" and "fraud" language, reddit and trustpilot cited
Best sustainable household brand?	(not asked)	Blueland, Grove, Seventh Gen, Dropps named. Public Goods absent.

ChatGPT's prior knowledge of Public Goods is relatively benign. The damage comes when the model actually reads the live web.

Every AI platform that matters moved to live web retrieval as default in 2026. The volume of queries that hit Pass 2 increases every month. Public Goods loses more ground every week without doing anything wrong. The clock is the cost.

Competitive position

Public Goods plus three closest competitors.

	Public Goods	Grove	Thrive Market	Blueland
Composite AI Visibility	52 (D)	81 (B+)	73 (B-)	80 (B+)
Platform	Shopify	Shopify	Custom	Shopify
UCP version	2026-04-08	2026-04-08	none	2026-04-08
UCP capabilities	8	8	0	8
llms.txt flavor	protocol template	protocol template	content map	content map
llms.txt size	4.4 KB	4.3 KB	396 KB	12.9 KB
JSON-LD SKU populated	no	yes	n/a	yes
Variant offers in schema	1 flat	per-variant	n/a	2 offers
FAQPage schema	no	no	n/a	yes
Wikipedia entity	none	yes	yes	none
Wikidata entity	none	Q111170467	Q27150214	none
Trustpilot	4.3 / 687	4.9 / 1,000+	4.3	4.5 / 812
Replies to negative reviews	no	yes	yes	yes
BBB rating	F / 9 unanswered	not audited	not audited	not audited
Named in ChatGPT "best brands"	no	yes	not asked	yes

Read

Public Goods is at parity with Grove on transactional infrastructure. That is the only axis where parity holds. On content depth, entity strength, recommendability, and reputation, Public Goods sits below every comparison brand including Blueland, which is a smaller company.

Thrive made the inverse bet: no UCP, no MCP, no agent-readable checkout. Instead they whitelisted 13 AI bots in robots.txt and built a 396 KB llms.txt indexing 2,785 categories. The cost asymmetry between Public Goods's llms.txt (4.4 KB) and Thrive's (396 KB) is 90 times. Thrive wins category recommendation queries against Public Goods because of it. Public Goods paid for Shopify-default infrastructure and skipped the content layer where the actual differentiation lives.

Proposed engagement

Two phases. Phase 1 is a 30-day stabilization sprint that closes the bleeding. Phase 2 is ongoing AI visibility operations that take the composite score from D to B+ over four to six months. After Phase 1 we measure against benchmarks set together and decide whether Phase 2 makes sense.

PHASE 1 · 30 DAYS

Stabilization sprint

Eight remediations sequenced as listed. Each addresses a finding above.

1. `Noindex /pages/legacy-membership`, remove from sitemap, request Google removal. Day 1.
2. Unify free shipping policy across FAQ, support, policy page, cart UI, and JSON-LD `OfferShippingDetails`. 301 retired URLs. Days 1 to 3.
3. Per-variant offers in JSON-LD across the catalog. Populate SKU and GTIN from existing product JSON. Days 3 to 8.
4. FAQPage schema across the 31 PDP accordion FAQs and the site-wide FAQ. Days 5 to 10.
5. Review schema alongside existing AggregateRating. Days 8 to 12.
6. Member-tier pricing as JSON-LD `UnitPriceSpecification` with `eligibleCustomerType`, integrated with Shopify customer-tag system. Populate empty `member-price` spans server-side. Days 10 to 20.
7. Customize `llms.txt` as content map: sustainability story, materials, founder, category structure, hero collections, certifications. Days 12 to 22.
8. Sitemap hygiene. Remove 12 to 14 administrative URLs. Days 22 to 26.

Days 26 to 30: verification, schema validator at 100%, benchmark re-measurement.

\$1,000 all-in for Phase 1.

Audit document, implementation, benchmark setup, verification. No separate audit fee.

Phase 1 benchmarks (finalized at kickoff)

- Composite AI Visibility Score (target: 52 to 65+, D to C).
- Frequency of the deprecated 40% offer in AI responses (target: 0 across 20 prompts).
- Coverage rate in ChatGPT, Perplexity, Gemini for the 20 audit prompts (before vs after).
- Schema validator pass rate across the catalog (target: 100%).
- Microsoft Copilot UCP poll frequency from server logs.

Phase 2 and beyond

Phase 1 closes the open wounds. Phase 2 takes Public Goods from D to B+. The work below is not what an in-house Shopify engineer does on a Tuesday.

Continuous AI behavior monitoring

Weekly 70-prompt benchmark across ChatGPT, Claude, Perplexity, Gemini, Copilot, Google AI Mode. Trend line per dimension, drift alerts, emerging-issue surfacing before retail metrics move. Internal teams do not have the prompt set, scoring discipline, or cross-platform access to build this.

Entity establishment campaign

Wikipedia article creation against the brand-name-disambiguation collision (the "Public Goods" string maps to the economics concept by default). Wikidata completion. Knowledge Panel population through structured Organization schema and sameAs linking. Editorial work requiring specific expertise in what AI engines treat as authoritative.

AI training corpus management

Reddit, YouTube, and community work calibrated to AI training data signals. Not generic PR. Identifying which threads AI engines surface, which language gets pulled verbatim, where counter-narrative work has leverage. BBB and Trustpilot response strategy connects here, written with awareness of how AI engines will read replies downstream.

Competitive intelligence

Tracking Grove, Thrive, Blueland, and emerging entrants. Schema changes, llms.txt expansions, certifications, Wikipedia edits, AI shopping platform partnerships. Most signals lead retail traffic by weeks. Anticipating them lets Public Goods move first.

AI shopping platform partnership development

ChatGPT Shopping, Perplexity Shop merchant list, Google AI Mode shopping integrations. Not self-serve. Require relationships, positioning, and product-data hygiene meeting platform-specific requirements. Public Goods is in none of them today.

Schema and protocol evolution

UCP advances by version. AP2 rolls out as a payment mandate standard. MCP server architecture is the new agent-store battleground. Staying current requires knowing what is coming, what matters, and what to ignore. Phase 2 pricing discussed after Phase 1 benchmarks land.

What only we do

The Phase 1 work could be approached by an internal team that knew it needed doing. The audit you are reading is what told the team it needed doing, and the methodology behind it is specifically what makes the difference.

We built the audit you are reading

Every finding came from running specific tools and queries in a specific order, with specific knowledge of what to look for. The audit took roughly forty hours across four people. None of it required admin access. All of it required knowing the right questions. Public Goods's in-house team is excellent at Shopify. We are excellent at AI visibility. The skill sets do not overlap as much as they look.

The 70-prompt benchmark infrastructure

Our evaluation set covers discovery, reputation, membership, shipping, category recommendation, comparative, attribute-specific, direct shopping intent, and hallucination probes. Scoring covers coverage, sentiment, citation source distribution, hallucination rate, and competitive share of voice. The benchmark is how we measure whether technical fixes are actually moving AI behavior. Internal teams typically need a year of trial and error to build something equivalent.

Cross-platform AI literacy

What works on ChatGPT does not work on Perplexity. What surfaces on Gemini is invisible on Claude. The category-recommendation behavior on Google AI Mode is different from the head-to-head comparison behavior in ChatGPT Shopping. We track these differences week to week.

Reddit and community signal calibration

Not all negative Reddit threads matter. Some are upvoted and indexed by AI engines as authoritative; others die in their subreddit. We identify which threads AI engines actually surface, which language they pull verbatim, and where counter-narrative work has leverage. This is research, not PR.

Schema engineering for emerging standards

JSON-LD `UnitPriceSpecification` with eligibility is not a standard Shopify pattern. UCP capability extensions for member-pricing-aware checkout are not in the Shopify docs. AP2 mandate integration is not a self-serve feature. These specific patterns are what take AI visibility from D to A.

The team

Armaan Priyadarshan (Dartmouth, Math and CS), **Samarjit Deshmukh** (Dartmouth, CS), **Ethan Zhou** (UC Berkeley, MET EECS and Business), **Shamit D'Souza** (Stanford, CS and Biology). Background spans AI agent research and infrastructure across 2025 to 2026, including custom MCP server work, agent evaluation frameworks, and AI visibility audits across consumer DTC categories.

Reputation surface, in detail

Every signal in the deck is publicly verifiable. This page indexes the source material AI engines and human researchers can both retrieve. Each item includes platform, identifier, and verbatim text where relevant.

Trustpilot, in full

4.3/5 average from 687 reviews. 74% five-star, per ChatGPT's own surfacing. Public Goods has the fewest Trustpilot reviews of the four brands compared. The Trustpilot profile carries a flag: *"Hasn't replied to negative reviews."*

1-star review titles, verbatim:

- *"Fraudulent 'membership' fees charged years after one purchase."* Body: single one-time purchase in 2020, charged \$79 in 2025, another \$79 in 2026.
- *"Fraudulent Cancellation Process."* Body: canceled in January, received email confirmation, charged again March 18.
- *"YOU CAN NEVER CANCEL!! Surreptitiously opted me into a membership."*
- *"Absolutely HORRENDOUS customer service... they gaslight you... I believe the 5 star reviews are not real."*

The word *Fraudulent* appears in two separate review titles. AI engines retrieving Trustpilot for citation pick up this exact language.

Reddit, by thread

Subreddit and thread	Upvotes	Date	Signal
r/Frugal · "Public Goods has been a Huge Letdown"	317	2019-03	Flagship critical thread; 58 comments
r/ZeroWaste · "PSA: Public Goods is a scam"	14	2024-08	Cited by ChatGPT alongside BBB
r/Scams · "Public Goods charged my card 4 years after purchase"	6	2024-08	Subreddit name itself is a signal AI engines weight
r/ZeroWaste · "Bad experience"	13	2021-04	Membership friction pattern
r/BuyItForLife · "Is Public Goods BIFL?"	0	2024-07	Ratio 0.27 (downvoted); ranks page 1 of Google for brand name
r/Frugal · "Has anyone tried"	6	2019-08	Mixed-to-negative replies
r/ZeroWaste · "Has anyone tried"	6	2020-05	Mixed-to-negative replies
r/Frugal · "2025 updates?"	1	2025-01	Sentiment continuing into 2025

Reddit top comments, verbatim

- 100up: *"the premise of 'at cost' is kind of illegitimate because every product is sold at an inflated rate."* Directly attacks the brand's core value proposition.
- 27up / 14up: *"they require a membership. eff that."* Pattern: free trial, order, cancel.
- 14up: *"It was a cash grab scam. I apologize for posting it."* An OP retracting a recommendation.
- 3up: *"they really do suppress negative reviews, which was a deal breaker... I found Grove to be a little better."*
- Multiple threads: *"plastics galore"* (greenwashing accusation, undermines sustainability positioning).

Listicle inclusion vs comparison exclusion

Public Goods appears in 8 of 40 third-party listicles we audited. Grove 8/40. Thrive 9/40. Blueland 4/40. Listicle presence is at parity. Where the brand disappears is head-to-head comparison content: the top "best sustainable brand vs" article on Google compares Grove, Blueland, CleanCult, Thrive, and Honest. Five brands, no Public Goods. ChatGPT pulls from comparison content disproportionately when synthesizing recommendations. Tier-1 press presence (CNN, USA Today, NPR, Forbes) is unconfirmed; those sites blocked our crawler.

Competitors named inside Public Goods Reddit threads

Thrive Market, Grove Collaborative, Blueland, The Good Store, SustainLA, Vitacost, Amazon. AI engines that follow citation chains end up at these brands when starting from a Public Goods search.

Technical surface and AI behavior, in detail

Product detail page JSON-LD inventory

Each Public Goods PDP carries 8 JSON-LD blocks: `Product` , `Organization` , `BreadcrumbList` , `ProductGroup` , `WebPage` , `Organization` (duplicate), `ProductGroup` (duplicate), `AggregateRating` . The `Product` node has `sku=null` and `mpn` populated with the Shopify product ID (8655369011338), not a GTIN. Sample PDPs: Natural Shampoo \$19.95 / 4.6 / 937 reviews. Laundry Sheets \$19.95 / 4.9 / 24. Tree-Free Toilet Paper \$49.95 / 4.4 / 473. Ceramic Bowls \$19.95 / 4.8 / 50. Medium Roast Coffee \$15.95 / 4.5 / 300.

Schema gap, by competitor

- **Grove Collaborative.** Per-product SKU populated (example: `MM-37-030`). Five per-variant offers per product, each with GTIN-12 (example: `808124144484`), availability, price, and `?variant=` deep link.
- **Blueland.** Kit-level SKU populated (example: `KP-BC-KIT-2P-EM`). `FAQPage` schema, `hasCertification` , `additionalProperty` . The smallest brand in the set carries the deepest schema.
- **Public Goods.** No SKU on any product, JSON-LD has no usable identifier. Real SKUs (`010-101-014-011`) and barcodes (`810049440107`) exist but live only in `/products/{handle}.json` , which AI engines do not consume as catalog truth. No `FAQPage` schema despite 31 HTML accordion mentions across the site. No Review schema. No `priceSpecification` .

Hidden product data and inadvertent exposure

The `/products/{handle}.json` endpoints expose information the JSON-LD layer should but does not carry: free-from claims (parabens, phosphates, chlorine bleach, added dyes, 1,4-dioxane) in `body_html` , "Made in Canada" origin, ply and sheet counts. The same endpoints also leak data the brand likely does not intend to expose: live inventory quantities (167, 91, etc.) and internal taxonomy tags (`__label: Best Seller` , `CLP2026-...` , `refills`). The valuable claims live where AI cannot read them; the operational data leaks where it should not.

Member pricing invisibility, three surfaces

Static HTML carries 12 `member-price` spans, all empty JavaScript placeholders.

`/products/{handle}.json` has no member or discount or tier keys; only a tag value "Member Discount" as a flag, with no numeric value. JSON-LD has no `priceSpecification` across any of the 8 blocks. The member discount is computed client-side and exposed in no machine-readable field. The brand's core value proposition is the one number AI cannot retrieve.

llms.txt and discovery layer, by competitor

Brand	llms.txt	llms-full.txt	Strategy	Sitemap loc
Public Goods	4.4KB, 86 lines	Near-identical to llms.txt	Shopify protocol template (how to transact, no catalog)	263
Grove	4.3KB	Near-identical to llms.txt	Same protocol template as Public Goods	7 (sitemap index)
Blueland	12.9KB	404	Content map (pages, collections, hacks with descriptions)	7 (sitemap index)
Thrive	395KB, 2,815 lines	2.7MB full catalog	Massive content map plus full product corpus	15,598

Thrive's `robots.txt` whitelists 13 named AI bots with `Crawl-delay: 5`. Public Goods and Grove have no AI-specific directives in `robots.txt`. Blueland allows OAI-SearchBot.

UCP and MCP, what we found

UCP discovery is live at `/well-known/ucp` at version 2026-04-08, all 8 `dev.ucp.shopping` capabilities exposed, both Google Pay (merchant ID 8387991) and Shopify Card payment handlers configured. The 13 canonical UCP Shopping methods are implemented. The MCP endpoint at `POST /api/ucp/mcp` is gated: agent profile must be supplied in `params.arguments.meta["ucp-agent"].profile` as a fetchable HTTPS URL with valid cache-control. Anonymous probes return 422 across the validation pipeline (missing profile URI, missing version, invalid cache control). The transactional infrastructure is correctly implemented and protocol-compliant.

The investment is invisible to ChatGPT. ChatGPT scrapes the page rather than transacting through MCP. From a live capture (P63, web search on): *"I can't directly add it to your cart, no access to your logged-in shopping session."* ChatGPT linked out, read the live page, surfaced the "Sold Out, Notify Me" state, and suggested an \$11.95 Organic Shampoo Bar as alternative. The UCP/MCP layer is well built and presently un-used by the dominant AI shopping channel.

ChatGPT behavior, captured

Accurate cart math (P62, web search). Computed Shampoo \$19.95 + Dish Soap \$9.95 = \$29.90, flagged \$15.10 short of free shipping at \$45. **Hallucinations against current ground truth.** P13 surfaced membership "often around \$79" (the current price is \$65; \$79 carried from older BBB complaint text). P17 surfaced stray "\$59/\$49" pricing (current is \$65). **Self-flagged contradiction (P02).** ChatGPT named all three shipping cohorts and flagged the contradiction unprompted: *"their current membership page ALSO says... no minimum required."*

Methodology. Crawler probes used curl with explicit AI bot user-agents (GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot) and Cloudflare returned HTTP 200 across all probes. Headless browser was bot-blocked on Google, DuckDuckGo, Trustpilot, and Reddit; headed Chromium with stealth cleared Google and Trustpilot. ChatGPT Plus accessed by importing valid session cookies into headed Chromium. Pass 1 (no browsing) and Pass 2 (forced web search) prompts captured across 40 prompts spanning discount, shipping, reputation, shopping, competitive, and category awareness queries. Raw evidence, screenshots, per-prompt response files, and TSV are available on request.

What is working, and next steps

What is working

Worth stating before closing. The transactional infrastructure Public Goods has built is sophisticated and worth preserving. UCP discovery is live at `/ .well-known/ucp` at version 2026-04-08 with all eight `dev.ucp.shopping` capabilities exposed. Both Google Pay and Shopify Card payment handlers are configured. The `llms.txt` and `agents.md` files are customized rather than Shopify default. The `robots.txt` is sophisticated and routes agents correctly with an explicit human-approval invariant for checkout. Bot accessibility is clean across all 32 probes we ran.

The work in this audit is not "rebuild Public Goods." It is "use the infrastructure you have already paid for, and add the content layer that makes it pay off." The brand is one or two quarters of focused work away from a B+.

Engagement form

Consulting. Not a license. Not a subscription. Not a platform sale. We deliver Phase 1 against benchmarks set together. After the sprint we measure against those benchmarks, look at the numbers together, and decide whether to continue into Phase 2.

Decision needed

Greenlight on Phase 1 (\$1,000) and a kickoff date. Benchmarks set in the first working session. Sprint begins within 5 business days of kickoff.

Contact

armaan.priyadarshan.29@dartmouth.edu

NOTE ON REPRODUCIBILITY

Every finding in this deck is externally reproducible using public tools (curl, Schema.org validator, ChatGPT Plus, Trustpilot, BBB, Reddit search, Shopify Commerce Readiness Scanner). We did not log into anything we did not have customer-tier access to. Methodology and raw data are available on request.