

BLACKWELL ENTERPRISES

Public Goods AI Visibility Audit

An external audit of how Public Goods performs across six AI platforms and our suggested course of action.

Prepared for Public Goods leadership

Date May 27, 2026

Contact armaan.priyadarshan.29@dartmouth.edu, samarjit.deshmukh.29@dartmouth.edu

Prepared by **Armaan Priyadarshan, Samarjit Deshmukh, Ethan Zhou, Shamit D'Souza**

Dartmouth Math & CS · Dartmouth CS · UC Berkeley MET (EECS & Business) · Stanford CS & Biology

Executive Summary

Our audit found three issues with how AI assistants surface products distributed by Public Goods.

1

Public Goods is absent from synthesized AI category recommendations.

When prompted for "best sustainable household brands," ChatGPT names Blueland, Grove, Seventh Generation, and Dropps. However, Public Goods is not mentioned at any depth. While the brand does appear in 8 of 40 human-written third-party listicles we audited, it disappears when AI engines synthesize their own rankings for several reasons. Public Goods has no Wikipedia entity or Wikidata, and its Knowledge Panel is visible only to merchants.

The Knowledge Panel for Public Goods carries delivery thresholds and payment options but zero entity facts (no founder, no founding date, no HQ). AI engines pulling "what is Public Goods" from the Knowledge Graph get a merchant card and nothing else.

2

AI engines tell prospects Public Goods is a scam.

AI assistants reference negative reviews when prompted about Public Goods. These include a BBB F rating, nine unanswered complaints, and verbatim "fraudulent membership" language from Trustpilot and Reddit. In addition, three independent platforms complain about unauthorized membership renewals years after a single purchase.

"BBB shows PublicGoods.com with an F rating, citing failure to respond to 9 complaints. Multiple 2025-2026 complaints about customers claiming they did not knowingly sign up for an annual membership." ChatGPT, web search on.

3

The free shipping policy contradicts itself across four pages.

Public Goods publishes conflicting free-shipping rules for three customer cohorts across four different crawlable pages. Members who joined after April 2025 get free shipping with no minimum, members who joined earlier need \$45, while non-members see a \$45 threshold on the FAQ page but \$65 on the support page. ChatGPT surfaces all three cohorts and flags the FAQ-versus-support contradiction to the user without being asked, which discourages prospects from checking out or building bigger baskets.

"The shipping FAQ on one page lists the \$45 threshold as universal for non-members, but the support page lists \$65. The current threshold depends on which page is treated as authoritative." ChatGPT, web search on.

Composite AI Visibility: **D (52/100)**, lowest in the competitive set, 27 points behind the average. Three further findings (member pricing invisibility, bait-and-switch from a deprecated offer, thin product schema) detailed after the scorecard.

AI Visibility Scorecard

We graded Public Goods on five dimensions that determine how AI engines treat brands. Each dimension was scored 0 to 100 against observed AI behavior across ChatGPT, Reddit, Trustpilot, BBB, and the brand’s own indexed pages. The composite score was calculated by averaging scores across the five dimensions.



Public Goods is the lowest-scoring brand amongst competitors, trailing the average by 27 points. The competitive average across Grove (B+), Thrive (B-), and Blueland (B+) is 79. Public Goods is at 52, the only sub-B grade in the category.

While the transactional infrastructure is strong, three of the five dimensions that drive AI behavior toward customers (quotability, recommendability, reputation) are failing.

Dimension	Public Goods	Grove	Thrive	Blueland	What it measures
Discoverability	 75	85	90	75	Entity presence, agent discovery files, crawler access
Quotability	 40	75	90	80	Schema depth, structured policies, accuracy in machine-readable form
Recommendability	 30	85	90	90	Inclusion in AI category recommendations, credentials, listicle presence
Transactability	 80	80	30	75	UCP, MCP, payment handlers, agent-readable pricing
Reputation	 35	80	65	80	Sentiment in AI training corpus, BBB, review aggregator presence
Composite	52 (D)	81 (B+)	73 (B-)	80 (B+)	

Grading scale: A 90+, B 73-89, C 60-72, D 50-59, F below 50.

The short version

Public Goods distributes high quality products with customer-friendly shipping fees and membership deals, but traditional AI engines fail to recognize this due to problems with the landing page. AI assistants do not meaningfully recommend the brand in category queries. Moreover, when asked about Public Goods directly, chatbots surface a fraud-adjacent reputation, a deprecated offer the brand discontinued last year, and a shipping policy that contradicts itself across four pages.

Here are our six main findings, ordered in terms of importance.

1. **Public Goods is absent from every AI-synthesized "best sustainable brands" ranking we ran.** ChatGPT with web search names Blueland, Grove, Seventh Generation, and Dropps as category leaders. Public Goods is not mentioned even though Blueland is a smaller company than Public Goods.
2. **BBB F rating with nine unanswered complaints surfaces in ChatGPT responses,** alongside corroborating "scam" and "fraud" language from Reddit and Trustpilot. Three independent platforms carry the same complaint about unauthorized membership renewals years after a single purchase. ChatGPT cites PublicGoods.com by domain when surfacing the BBB record.
3. **Chatbots don't understand the free shipping policy.** Our audit found the landing page pulls from two contradicting sources across three customer cohorts, confusing AI assistants about the shipping policy. This disincentivizes users from buying products or building bigger baskets.
4. **Member pricing is structurally invisible across every machine-readable surface.** The membership's actual value proposition cannot be quoted by any AI agent that does not render JavaScript. Public Goods loses every head-to-head comparison this way. ChatGPT frames the brand as the pricier option versus Grove and Thrive because the discount cannot be quantified.
5. **A deprecated 40% membership offer is actively indexed and surfaces in ChatGPT responses** to prospective new members. The page calls 40% the "best available discount," but that offer was discontinued in April 2025. AI engines are bait-and-switching potential customers with the brand's own indexed content.
6. **The PDP schema is materially thinner than Grove's and Blueland's.** Product pages are missing per-variant offers and FAQPage schemas.

NOTE ON REPRODUCIBILITY

Every finding in this deck is externally reproducible using public tools (curl, Schema.org validator, ChatGPT Plus, Trustpilot, BBB, Reddit search, Shopify Commerce Readiness Scanner). We did not log into anything we did not have customer-tier access to. Our methodology and raw data are available on request.

Public Goods does not exist in AI category awareness

Prompt: *"Best sustainable household brands."* (ChatGPT with web search)

ChatGPT named four brands as category leaders, each with citable credentials:

- Blueland (best overall, citing B-Corp certification)
- Grove Collaborative (best one-stop, citing range and certifications)
- Seventh Generation
- Dropps

Public Goods was not in the main list or in any "consider also" section. The brand does not appear at any depth of the response.

Low visibility is expensive for Public Goods

Public Goods appears in 8 of 40 human-written third-party listicles we audited, but it does not appear meaningfully in AI recommendations. As AI shopping recommendations replace human listicles, Public Goods loses a significant source of inbound customers.

The brands ChatGPT named all have something to cite. For instance, Blueland and Grove have B-Corp certification, while Thrive (Q27150214) has Wikidata entities. The brand name also collides with the economics concept "public good," which dominates the SERP: four of the top eight organic Google results for "Public Goods" are the econ concept (Wikipedia "Public good", econlib, Investopedia, Khan Academy).

Three further signals AI engines weight

Merchant-only Knowledge Panel. Google's Knowledge Panel for Public Goods exists, but it's a commerce/merchant card (free delivery threshold, returns policy, payment options). Zero entity facts: no founder (Morgan Hirsh, founded 2016), no HQ, no founding date, no CEO. AI engines pulling "what is Public Goods" from the Knowledge Graph get nothing about the brand.

Negative Reddit ranks page 1 of Google for the brand's own name. Search "publicgoods.com": position 7 is the r/ZeroWaste "PSA: Public Goods is a scam" thread. Search "Public Goods brand": position 3 is the r/BuyItForLife "Is Public Goods BIFL?" thread (upvote ratio 0.27). AI engines that crawl SERP for citations pick these up.

Absent from head-to-head comparison content. The top-ranking "best sustainable brand vs" article on Google is titled "Grove vs Blueland vs CleanCult vs Thrive vs Honest." Five brands compared in the category, none of them Public Goods. This is what ChatGPT pulls from when synthesizing comparison answers. Closing the gap is editorial and entity work and can't be replicated with pre-built Shopify infrastructure.

AI engines tell prospects Public Goods is a scam

The screenshot displays the PublicGoods.com website. The top navigation bar includes links for 'Visit Website', 'Email Business', and 'Write a Review'. The main content area is divided into several sections:

- BUSINESS PROFILE:** Online Retailer, PublicGoods.com, Headquarters. A warning states: 'This business is NOT BBB Accredited. Find BBB Accredited Businesses in Online Retailer.'
- Overview:** Headquarters, PublicGoods.com, PO Box 838, New York, NY 10008-0838. A link to 'Find Local Branches' is provided.
- BBB Accreditation & Rating:** PublicGoods.com is NOT a BBB Accredited Business. To become accredited, a business must agree to BBB Standards for Trust and pass BBB's vetting process. A link to 'Why choose a BBB Accredited Business?' is provided.
- BBB Rating:** F. Reasons for rating: Failure to respond to 9 complaint(s) filed against business; 4 complaint(s) filed against business that were not resolved. A link to 'More Information' is provided.
- Latest Reviews:** A section titled 'How are BBB ratings calculated?' with a link to 'View all Reviews'.
- Join Trusted Businesses:** A button to 'Become BBB Accredited'.

F rating. NOT BBB Accredited. 9 unanswered complaints with 4 unresolved. File opened 5/29/2020. Reviewer language: "DO NOT TRUST", "complete scam", "fraudulent annual subscription."

Three independent platforms with the same complaint and emphatic language

"Wife purchased items 3-4 years ago. We have been charged \$65 annually for an 'annual subscription' service. (BBB)"

"Fraudulent membership fees charged years after one purchase. In 2025 I was charged \$79 for a membership I never authorized. (Trustpilot, 1-star)"

"my bank caught it and flagged it as fraud. (Reddit r/ZeroWaste, in a thread titled "PSA: Public Goods is a scam")"

Additional Trustpilot 1-star titles, verbatim: "Fraudulent Cancellation Process" (canceled in January, charged March 18). "YOU CAN NEVER CANCEL!! Surreptitiously opted me into a membership." "Absolutely HORRENDOUS customer service... they gaslight you.... I believe the 5 star reviews are not real." The word "Fraudulent" appears in two separate review titles. Reddit r/Scams also carries a 2024 thread: "Public Goods charged my card 4 years after purchase."

What ChatGPT tells a prospect today

Prompt with web search: "Public Goods customer complaints." Sources cited: bbb.org, trustpilot.com, reddit.com, publicgoods.com.

"BBB shows PublicGoods.com with an F rating, citing failure to respond to 9 complaints. Multiple 2025-2026 complaints about customers claiming they did not knowingly sign up for an annual membership."

The shipping policy contradicts itself across four pages

The shipping policy varies across three customer cohorts and two source pages. There are six possible answers that are all crawlable, causing significant confusion.

Customer cohort	Free shipping rule	Source page
Members enrolled after 4/1/2025	Free on all orders, no minimum	Membership info
Members enrolled before 4/1/2025	Free on orders of \$45+	Legacy membership
Non-members	\$45 minimum	FAQ page
Non-members	\$65 minimum	Support page

ChatGPT flags the contradiction unprompted

Prompt with web search: *"Does Public Goods offer free shipping?"*

"Their current membership page also says members get free U.S. shipping on every order with no minimum required. The shipping FAQ on a different page lists the \$45 threshold as universal for non-members, but the support page lists \$65. The current threshold depends on which page is treated as authoritative."

The policy is internally inconsistent enough that an AI engine flags it as a contradiction to the user. The user does not have to notice. ChatGPT volunteers the contradiction as part of the response.

Remediation requires more than picking one page and editing it. The policy must be unified across the FAQ, support, policy page, cart drawer, schema OfferShippingDetails, and the contradicting pages must be 301-redirected because AI engines have already cached the old versions.

The membership's actual value is invisible to AI

The entire reason to buy the membership is the lower price. That number exists in no machine-readable field anywhere on the site.

All three surfaces are empty

JSON-LD on the natural shampoo PDP. There's one flat Offer at \$19.95, but the priceSpecification, member tier, eligibility logic, SKUs and GTINs are missing.

The /products/natural-shampoo.json endpoint. The endpoint exposes a "Member Discount" tag as a string but no numeric discount.

Static HTML on the PDP. The twelve member-price spans are all empty. The actual member price is computed and injected by JavaScript at render time.

The competitive consequence we measured

Prompt: *"Public Goods vs Grove Collaborative."* ChatGPT with web search.

ChatGPT frames Public Goods as the more expensive membership: \$65/year against Grove's \$29.99/year. The 20% discount that would offset this cannot be quantified because the member price is invisible. Public Goods loses every head-to-head this way. Same dynamic with Thrive (\$59.95/year).

THIS IS NOT A GENERIC SHOPIFY THEME ISSUE

Member pricing as eligible-buyer schema requires custom UnitPriceSpecification with eligibleCustomerType, integrated with the Shopify customer-tag system, populated in both JSON-LD and the product JSON, and rendered server-side rather than client-side. No Shopify app does this out of the box for membership stores. The pattern requires specific schema knowledge and bespoke implementation work.

AI is bait-and-switching the brand's own prospects

The page at `/pages/legacy-membership` contains this sentence in its static, server-rendered HTML:

"automatically enrolled in our exclusive VIP Member program. This gives you the best available discount of 40% off most categories (30% off food items, regular pricing on frozen fish and plants), a special benefit..."

"Best available discount" implies the 40% rate is currently the best on offer but it is not. The April 2025 restructure moved new members to a flat 20%. Legacy members were grandfathered into the 40%. Nobody can sign up for that rate today.

Three things compound the problem:

- The page has no `noindex` meta directive or `X-Robots-Tag` header.
- The canonical URL is self-referencing. Public Goods is declaring the page the canonical version to index.
- It is present in `sitemap_pages_1.xml`. The site actively submits it to Google.

What ChatGPT does with it

Prompt with web search: *"What discount do Public Goods members get?"*

"20% off for new members. One caveat: Public Goods says legacy VIP members may keep a higher discount of 40% off most categories, 30% off food, and regular pricing on frozen fish and plants."

Source cited: `publicgoods.com`

ChatGPT presents the deprecated offer as a live "caveat" to a new prospect. The prospect signs up, gets the 20%, and learns the 40% rate was never available to them. The brand's own indexed content is producing the disappointment.

WHY THIS GETS WORSE OVER TIME

Every additional AI engine that adds live web retrieval is another engine pulling this page. The page is more useful to AI engines than to humans. Traffic to it is increasingly bots, which means its downstream impact is increasingly on AI responses, not on real customers seeing it directly.

Product detail schema is materially thinner than Grove and Blueland

Each Public Goods PDP carries 8 JSON-LD blocks (Product, Organization x2, BreadcrumbList, ProductGroup x2, WebPage, AggregateRating). The Product node has aggregateRating, brand, description, image, mpn, name, offers. What it does not have, by competitor benchmark:

- **No SKU.** sku field is null. Grove emits real SKUs (example: MM-37-030). Blueland too. Public Goods is the only one of the three with no SKU.
- **No per-variant offers.** One flat Offer at \$19.95 collapses real Travel/Full/Refill/Set variants ranging \$9.95 to \$49.95. Grove emits five per-variant offers, each with its own price, GTIN-12, availability, and ?variant= deep link URL. ChatGPT cannot tell the agent which variant to buy.
- **mpn holds the Shopify product ID, not a GTIN.** Real barcodes (example: 810049440107) exist in /products/{handle}.json but never make it into JSON-LD. No usable identifier reaches AI catalog graphs.
- **No FAQPage schema.** The PDP carries 31 accordion FAQ mentions in HTML. Blueland exposes FAQPage. Public Goods does not. Schema.org Question and Answer nodes are how AI engines source PDP-level Q&A.

Blueland additionally exposes hasCertification and additionalProperty for kit composition. Public Goods exposes neither. The smallest brand in the competitive set carries the deepest product schema.

The trajectory: AI behavior is shifting against the brand

We ran every prompt twice: once with web browsing disabled (Pass 1, training-data-only), once with web browsing enabled (Pass 2, live retrieval). The gap between the two passes is the gap that AI engines close every week as live retrieval becomes default behavior across ChatGPT, Perplexity, and Google AI Mode.

Prompt category	Pass 1: no browsing	Pass 2: web search on
Discount question	Says 20%, the correct current rate.	Adds 40% as a current "caveat" by pulling the deprecated page.
Reputation question	"Legitimate, not a scam, but watch the membership terms."	Surfaces BBB F rating, 9 unanswered complaints, "fraudulent" verbatim.
Shipping question	Flat \$45 minimum.	All three cohorts surfaced plus the FAQ-vs-Support contradiction flagged.
Category recommendation	Public Goods sometimes named alongside Grove and Thrive.	Public Goods absent. Blueland, Grove, Seventh Generation, Dropps named.

In 2024 a model’s training data was Public Goods’s primary AI exposure. In 2026 it is live retrieval, and the live retrieval surface is materially worse than the training data was.

Competitive position

This matrix details where Public Goods sits on the specific signals AI engines weight when deciding what to recommend and how to characterize a brand.

Signal	Public Goods	Grove	Thrive	Blueland
Wikipedia / Wikidata entity	No	Yes	Yes	No
B-Corp or comparable credential surfaced	No	Yes	Implicit	Yes
JSON-LD SKU on products	No	Yes	N/A	Yes
Per-variant offers in JSON-LD	No	Yes	N/A	Yes
FAQPage schema	No	No	No	Yes
llms.txt strategy	Protocol template	Protocol template	Massive content map	Content map
UCP / MCP discovery	Live	Live	Not present	Live
Trustpilot score	4.3 / 687	4.9 / 1K+	4.3	4.5
BBB status	F, not accredited	Not flagged	Not flagged	Not flagged
Composite AI Visibility	52 (D)	81 (B+)	73 (B-)	80 (B+)

This matrix implies three important things. First, Public Goods's transactional layer is at parity with Grove and Blueland and ahead of Thrive. Second, Public Goods is falling behind on almost every signal that drives AI recommendation behavior (entity, credentials, schema depth, llms.txt content). Third, the BBB F rating is a category outlier with Grove (4.9), Blueland (4.5), and Thrive (4.3) all maintaining a level of trust that Public Goods does not.

Proposed engagement

Our proposed engagement would occur in two phases. Phase 1 is a 30-day sprint to fix the glaring issues with Public Goods' AI visibility. Phase 2 involves ongoing AI visibility operations that take the composite score from D to A- over four to six months. After Phase 1, we will measure against benchmarks set together and decide whether Phase 2 makes sense.

PHASE 1 · 30-DAY STABILIZATION SPRINT

We propose seven specific remediations to close findings identified in this audit.

1. **noindex the deprecated 40% legacy-membership page, remove from sitemap.** This is 1 line of code and a sitemap re-submit that ensures LLMs are not confused about membership pricing.
2. **Unify the shipping policy across the FAQ, support page, membership info, and policy page.** Add OfferShippingDetails schema and 301-redirect the contradicting versions.
3. **Implement custom UnitPriceSpecification with eligibleCustomerType for member pricing.** Populate in JSON-LD and product JSON and start rendering server-side.
4. **Add per-variant Offer nodes with real SKUs and GTINs to PDP JSON-LD.** We would match the per-variant pattern Grove uses.
5. **Add FAQPage schema to PDPs that have accordion FAQs (31 pages identified).** This closes the FAQPage gap in Finding 6.
6. **Add Review schema (not just AggregateRating) to PDPs.** Pulls in individual customer reviews as Schema.org Review nodes.
7. **Custom llms.txt content map with brand context, materials, and category structure.** Replace the Shopify protocol template with a content-first map matching Blueland's pattern.

Benchmarks set at kickoff

We propose the following joint benchmarks measured before and after the sprint.

- Composite AI Visibility Score (target: 52 to 60+, D to C).
- Schema validator pass rate across PDPs (target: 100%).
- Surfacing rate of the deprecated 40% offer in ChatGPT (target: 0% from any new-prospect query).
- Shipping policy contradiction surfacing rate (target: 0% across 10 phrasings of the shipping prompt).
- Coverage rate in ChatGPT, Perplexity, Gemini "best sustainable household" responses (target: appearance in at least 2 of 6 prompts where 0 appear today).

PHASE 1 FEE

\$1,000 USD for the 30-day engagement in accordance to the benchmarks set above. Payable on delivery of the verification report at day 30.

Phase 2 and beyond

During Phase 2, we will help Public Goods improve its AI visibility further over the following months. In particular, this includes the following:

Continuous AI behavior monitoring

Weekly 70-prompt benchmarks across ChatGPT, Claude, Perplexity, Gemini, Copilot, Google AI Mode. Includes trend line per dimension, drift alerts, and emerging-issue surfacing before retail metrics move.

Entity establishment campaign

Wikipedia article (drafted to notability standards), Wikidata entry with disambiguation against the econ concept Q1673101, Knowledge Panel facts (founder, founding date, HQ, CEO, B-Corp if pursued). About 8 to 12 weeks of effort.

AI training corpus management

Reddit, YouTube, and community work calibrated to AI training data signals. Identifying which threads AI engines surface, which language gets pulled verbatim, where counter-narrative work has impact. The BBB and Trustpilot response strategy connects here, written with awareness of how AI engines will read replies downstream.

Competitive intelligence on agent-readable surfaces

Weekly tracking of what Grove, Thrive, and Blueland ship on their llms.txt, agents.md, JSON-LD, and protocol surfaces. Implement counter-moves in areas of high leverage.

AI shopping platform partnership development

ChatGPT Shopping (announced expansion), Perplexity Shop, Google AI Mode shopping integrations. These platforms onboard partners through specific schema and policy requirements that go beyond what is currently public.

Schema and protocol engineering for emerging standards

JSON-LD UnitPriceSpecification with eligibility is not a standard Shopify pattern. UCP capability extensions for member-pricing-aware checkout are not in the Shopify docs. AP2 mandate integration is not a self-serve feature. These specific patterns are what take AI visibility from D to A.

What only we do

The team

Armaan Priyadarshan (Dartmouth, Math and CS), **Samarjit Deshmukh** (Dartmouth, CS), **Ethan Zhou** (UC Berkeley, MET EECS and Business), **Shamit D'Souza** (Stanford, CS and Biology).

We're a highly technical team with significant DTC domain expertise, working with 10+ brands valued over \$300M. Our backgrounds span AI agent research and infrastructure across 2025 to 2026, including custom MCP server work, agent evaluation frameworks, and AI visibility audits across consumer DTC categories.

We built the audit you are reading

Every finding came from running specific tools and queries in a specific order, with specific knowledge of what to look for. This audit took roughly thirty hours across four people. In-house Shopify tools provide some insight, but only we have the expertise needed to build quality AI visibility infrastructure.

The 70-prompt benchmark infrastructure

Our evaluation set covers discovery, reputation, membership, shipping, category recommendation, comparative, attribute-specific, direct shopping intent, and hallucination probes. Scoring covers coverage, sentiment, citation source distribution, hallucination rate, and competitive share of voice. This benchmark is how we measure whether technical fixes are actually moving AI behavior. Internal teams typically need a year of trial and error to build something equivalent.

Cross-platform AI literacy

What works on ChatGPT does not work on Perplexity. What surfaces on Gemini is invisible on Claude. The category-recommendation behavior on Google AI Mode is different from the head-to-head comparison behavior in ChatGPT Shopping. We spend time tracking these differences week to week so that you don't have to.

Reddit and community signal calibration

Not all negative Reddit threads matter. Some are upvoted and indexed by AI engines as authoritative, while others die in their subreddit. We identify which threads AI engines actually surface, which language they pull verbatim, and where counter-narrative work has impact.

Schema engineering for emerging standards

Shopify doesn't provide out-of-the-box solutions for AI visibility. For instance, JSON-LD UnitPriceSpecification with eligibility is not a standard Shopify pattern, UCP capability extensions for member-pricing-aware checkout are not in the Shopify docs, and AP2 mandate integration is not a self-serve feature.

Reputation surface, in detail

Every signal in the deck is publicly verifiable. This page indexes the source material AI engines and human researchers can both retrieve. Each item includes platform, identifier, and verbatim text where relevant.

Trustpilot, in full

4.3/5 average from 687 reviews. 74% five-star, per ChatGPT's own surfacing. Public Goods has the fewest Trustpilot reviews of the four brands compared. The Trustpilot profile carries a flag: "Hasn't replied to negative reviews."

1-star review titles, verbatim:

- "Fraudulent 'membership' fees charged years after one purchase." Body: single one-time purchase in 2020, charged \$79 in 2025, another \$79 in 2026.
- "Fraudulent Cancellation Process." Body: canceled in January, received email confirmation, charged again March 18.
- "YOU CAN NEVER CANCEL!! Surreptitiously opted me into a membership."
- "Absolutely HORRENDOUS customer service... they gaslight you... I believe the 5 star reviews are not real."

The word "Fraudulent" appears in two separate review titles. AI engines retrieving Trustpilot for citation pick up this exact language.

Reddit, by thread

Subreddit and thread	Upvotes	Date	Signal
r/Frugal · "Public Goods has been a Huge Letdown"	317	2019-03	Flagship critical thread; 58 comments
r/ZeroWaste · "PSA: Public Goods is a scam"	14	2024-08	Cited by ChatGPT alongside BBB
r/Scams · "Public Goods charged my card 4 years after purchase"	6	2024-08	Subreddit name itself is a signal AI engines weight
r/ZeroWaste · "Bad experience"	13	2021-04	Membership friction pattern
r/BuyItForLife · "Is Public Goods BIFL?"	0	2024-07	Ratio 0.27 (downvoted); ranks page 1 of Google for brand name
r/Frugal · "Has anyone tried"	6	2019-08	Mixed-to-negative replies
r/ZeroWaste · "Has anyone tried"	6	2020-05	Mixed-to-negative replies
r/Frugal · "2025 updates?"	1	2025-01	Sentiment continuing into 2025

Reddit top comments, verbatim

- 100up: "the premise of 'at cost' is kind of illegitimate because every product is sold at an inflated rate." Directly attacks the brand's core value proposition.
- 27up / 14up: "they require a membership. eff that." Pattern: free trial, order, cancel.
- 14up: "It was a cash grab scam. I apologize for posting it." An OP retracting a recommendation.
- 3up: "they really do suppress negative reviews, which was a deal breaker... I found Grove to be a little better."
- Multiple threads: "plastics galore" (greenwashing accusation, undermines sustainability positioning).

Listicle inclusion vs comparison exclusion

Public Goods appears in 8 of 40 third-party listicles we audited. Grove 8/40. Thrive 9/40. Blueland 4/40. Listicle presence is at parity. Where the brand disappears is head-to-head comparison content: the top "best sustainable brand vs" article on Google compares Grove, Blueland, CleanCult, Thrive, and Honest. Five brands, no Public Goods. ChatGPT pulls from comparison content disproportionately when synthesizing recommendations. Tier-1 press presence (CNN, USA Today, NPR, Forbes) is unconfirmed; those sites blocked our crawler.

Competitors named inside Public Goods Reddit threads

Thrive Market, Grove Collaborative, Blueland, The Good Store, SustainLA, Vitacost, Amazon. AI engines that follow citation chains end up at these brands when starting from a Public Goods search.

Technical surface and AI behavior, in detail

Product detail page JSON-LD inventory

Each Public Goods PDP carries 8 JSON-LD blocks: Product, Organization, BreadcrumbList, ProductGroup, WebPage, Organization (duplicate), ProductGroup (duplicate), AggregateRating. The Product node has sku=null and mpn populated with the Shopify product ID (8655369011338), not a GTIN. Sample PDPs: Natural Shampoo \$19.95 / 4.6 / 937 reviews. Laundry Sheets \$19.95 / 4.9 / 24. Tree-Free Toilet Paper \$49.95 / 4.4 / 473. Ceramic Bowls \$19.95 / 4.8 / 50. Medium Roast Coffee \$15.95 / 4.5 / 300.

Schema gap, by competitor

- **Grove Collaborative.** Per-product SKU populated (example: MM-37-030). Five per-variant offers per product, each with GTIN-12 (example: 808124144484), availability, price, and ?variant= deep link.
- **Blueland.** Kit-level SKU populated (example: KP-BC-KIT-2P-EM). FAQPage schema, hasCertification, additionalProperty. The smallest brand in the set carries the deepest schema.
- **Public Goods.** No SKU on any product, JSON-LD has no usable identifier. Real SKUs (010-101-014-011) and barcodes (810049440107) exist but live only in /products/{handle}.json, which AI engines do not consume as catalog truth. No FAQPage schema despite 31 HTML accordion mentions across the site. No Review schema. No priceSpecification.

Hidden product data and inadvertent exposure

The /products/{handle}.json endpoints expose information the JSON-LD layer should but does not carry: free-from claims (parabens, phosphates, chlorine bleach, added dyes, 1,4-dioxane) in body_html, "Made in Canada" origin, ply and sheet counts. The same endpoints also leak data the brand likely does not intend to expose: live inventory quantities (167, 91, etc.) and internal taxonomy tags (__label: Best Seller, CLP2026-..., refills). Fundamentally, the valuable claims live where AI cannot read them.

Member pricing invisibility, three surfaces

Static HTML carries 12 member-price spans, all empty JavaScript placeholders. /products/{handle}.json has no member or discount or tier keys; only a tag value "Member Discount" as a flag, with no numeric value. JSON-LD has no priceSpecification across any of the 8 blocks. The member discount is computed client-side and exposed in no machine-readable field. The brand's core value proposition is the one number AI cannot retrieve.

llms.txt and discovery layer, by competitor

Brand	llms.txt	llms-full.txt	Strategy	Sitemap loc
Public Goods	4.4KB, 86 lines	Near-identical to llms.txt	Shopify protocol template (how to transact, no catalog)	263
Grove	4.3KB	Near-identical to llms.txt	Same protocol template as Public Goods	7 (sitemap index)
Blueland	12.9KB	404	Content map (pages, collections, hacks with descriptions)	7 (sitemap index)
Thrive	395KB, 2,815 lines	2.7MB full catalog	Massive content map plus full product corpus	15,598

Thrive's robots.txt whitelists 13 named AI bots with Crawl-delay: 5. Public Goods and Grove have no AI-specific directives in robots.txt. Blueland allows OAI-SearchBot.

UCP and MCP, what we found

UCP discovery is live at /.well-known/ucp at version 2026-04-08, all 8 dev.ucp.shopping capabilities exposed, both Google Pay (merchant ID 8387991) and Shopify Card payment handlers configured. The 13 canonical UCP Shopping methods are implemented. The MCP endpoint at POST /api/ucp/mcp is gated: agent profile must be supplied in params.arguments.meta["ucp-agent"].profile as a fetchable HTTPS URL with valid cache-control. Anonymous probes return 422 across the validation pipeline (missing profile URI, missing version, invalid cache control). The transactional infrastructure is correctly implemented and protocol-compliant.

The investment is invisible to ChatGPT. ChatGPT scrapes the page rather than transacting through MCP. From a live capture (P63, web search on): *"I can't directly add it to your cart, no access to your logged-in shopping session."* ChatGPT linked out, read the live page, surfaced the "Sold Out, Notify Me" state, and suggested an \$11.95 Organic Shampoo Bar as an alternative. The UCP/MCP layer is well built and presently un-used by the dominant AI shopping channel.

ChatGPT behavior, captured

Accurate cart math (P62, web search). Computed Shampoo \$19.95 + Dish Soap \$9.95 = \$29.90, flagged \$15.10 short of free shipping at \$45. **Hallucinations against current ground truth.** P13 surfaced membership "often around \$79" (the current price is \$65; \$79 carried from older BBB complaint text). P17 surfaced stray "\$59/\$49" pricing (current is \$65). **Self-flagged contradiction (P02).** ChatGPT named all three shipping cohorts and flagged the contradiction unprompted: *"their current membership page ALSO says... no minimum required."*

Methodology. Crawler probes used curl with explicit AI bot user-agents (GPTBot, ClaudeBot, PerplexityBot, OAI-SearchBot) and Cloudflare returned HTTP 200 across all probes. Headless browser was bot-blocked on Google, DuckDuckGo, Trustpilot, and Reddit; headed Chromium with stealth cleared Google and Trustpilot. ChatGPT Plus accessed by importing valid session cookies into headed Chromium. Pass 1 (no browsing) and Pass 2 (forced web search) prompts captured across 40 prompts spanning discount, shipping, reputation, shopping, competitive, and category awareness queries. Raw evidence, screenshots, per-prompt response files, and TSV are available on request.

What is working, and next steps

What is working

Our audit also found several areas where the Public Goods site is well maintained. In particular, the transactional infrastructure Public Goods has built is sophisticated and worth preserving. UCP discovery is live at /.well-known/ucp at version 2026-04-08 with all eight dev.ucp.shopping capabilities exposed, and both Google Pay and Shopify Card payment handlers are configured. The llms.txt and agents.md files are customized rather than Shopify default. The robots.txt is sophisticated and routes agents correctly with an explicit human-approval invariant for checkout. Overall, bot accessibility is clean across all 32 probes we ran.

Decision needed

Greenlight on Phase 1 (\$1,000) and a kickoff date. Benchmarks have been outlined above, and will be finalized before the first working session. We can provide a formal engagement letter if needed.

Contact

armaan.priyadarshan.29@dartmouth.edu, samarjit.deshmukh.29@dartmouth.edu